

# POSTGRES ON OPENSHIFT WORKSHOP



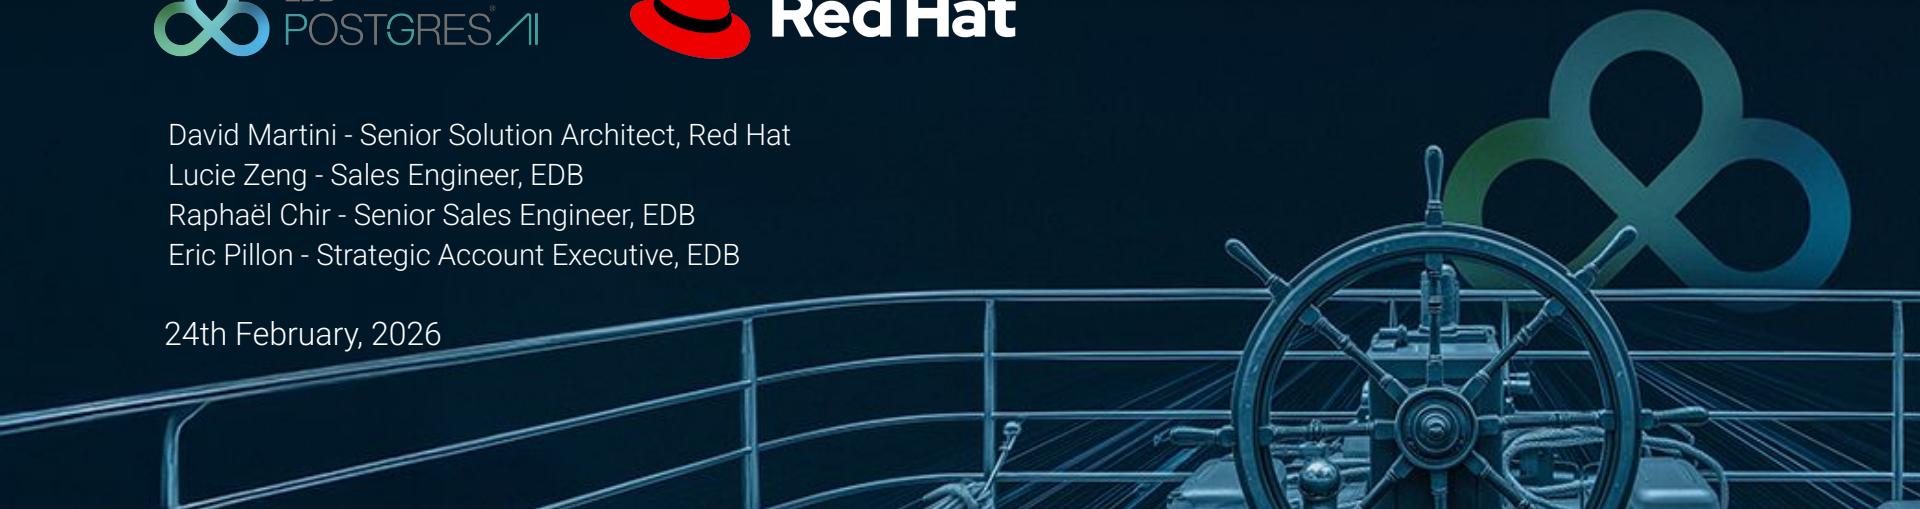
David Martini - Senior Solution Architect, Red Hat

Lucie Zeng - Sales Engineer, EDB

Raphaël Chir - Senior Sales Engineer, EDB

Eric Pillon - Strategic Account Executive, EDB

24th February, 2026



# Agenda

Commencer	Fin	Session
12:30	13:30	Accueil Cocktail Lunch
13:30	13:45	Partenariat entre Red Hat OpenShift et EDB (Red Hat - David Martini)
13:45	14:00	Présentation du marché Postgres et d'EDB (EDB - Eric Pillon)
14:00	14:30	Architecture de référence de l'opérateur CNPG (EDB - Raphaël Chir)
14:30	16:30	Atelier pratique (EDB - Lucie Zeng & Raphaël Chir)
16:30	17:00	Pause
17:00	17:30	Conclusion et perspectives



# Red Hat OpenShift with EDB

David Martini  
Senior Solution Architect

# The world's leading provider of open source enterprise IT solutions

More than  
**90%**  
of the  
Fortune  
**500**  
use  
**Red Hat**  
products and  
solutions<sup>1</sup>

**~22,000**  
employees

**105 +**  
offices

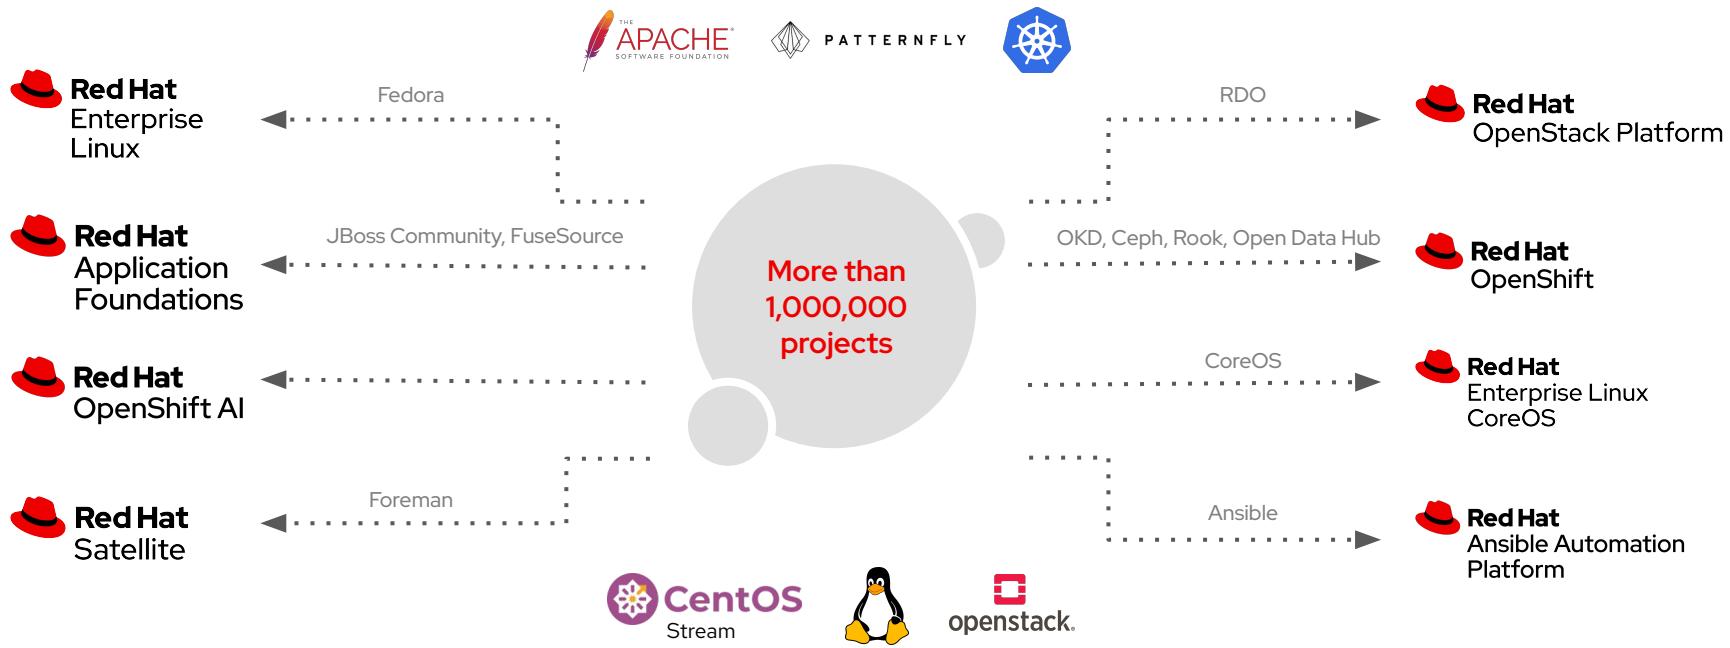
The first  
**\$3**  
**billion**  
open  
source  
company  
in the world<sup>2</sup>

# Open source

Open source is about more than developing software.  
It's how we built our company.

And it's why we have been so successful.

# From communities to enterprise



## Linux contributions



Platinum members



Red Hat has been working on Linux since 1993



Linux is the foundation of all Red Hat products



The first Patent promise was issued by Red Hat in 2002

## OpenInfra contributions



Historical member



First Red Hat OpenStack release in 2013



Most represented member on the board (3/21)



Red Hat & IBM = more than 3 times contributions as the second contributor

## Cloud-native contributions



Co-founder in 2015



Platinum members

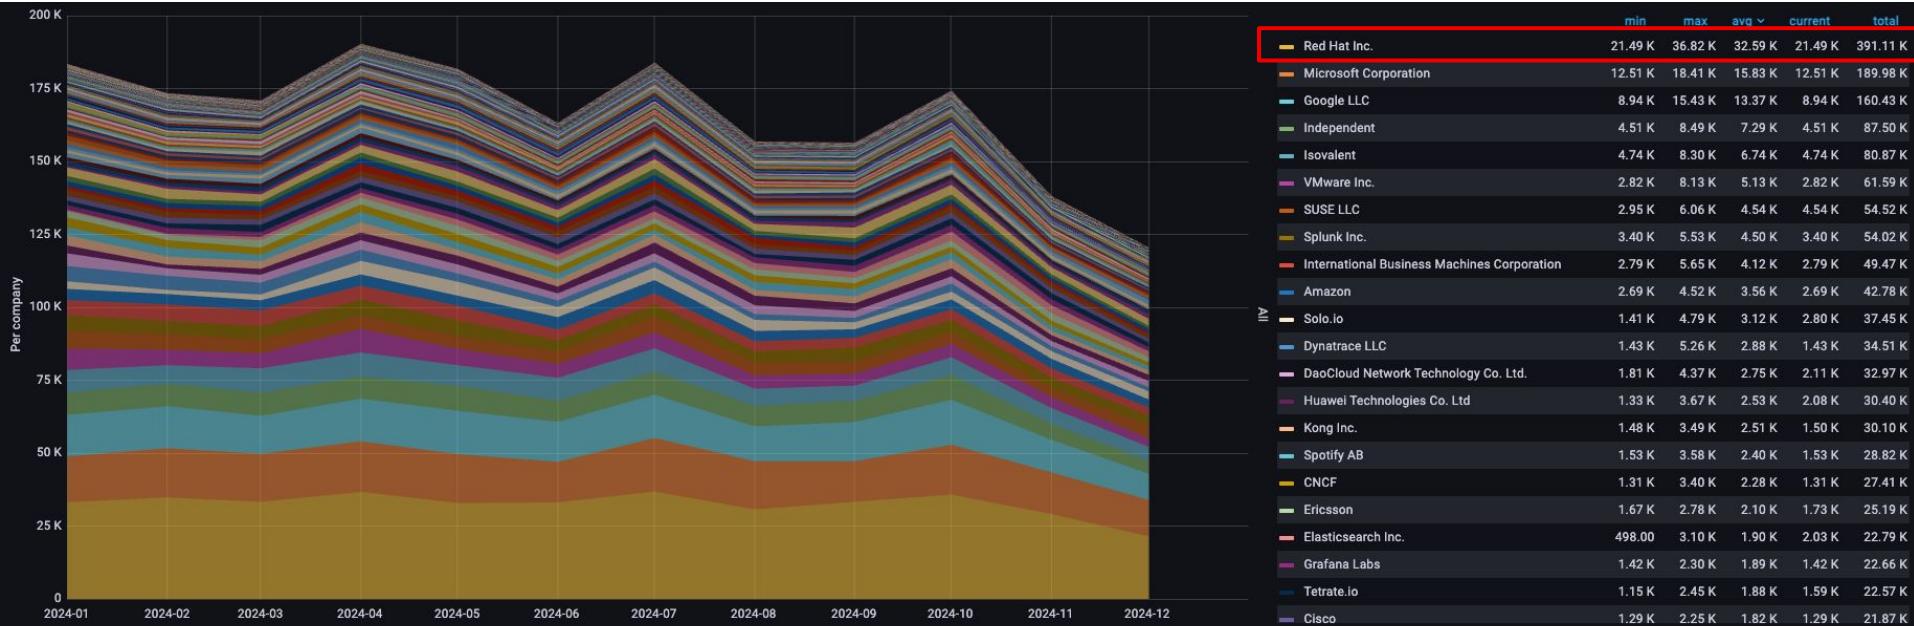


Second Kubernetes contributor after Google



First CNCF contributor in 2022, 2023, 2024, etc...

# CNCF 2024 Global contribution



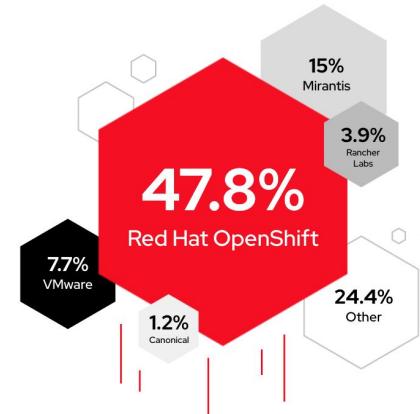
*Red Hat is the #1 company for contributions to CNCF in 2024.*

## Red Hat portfolio

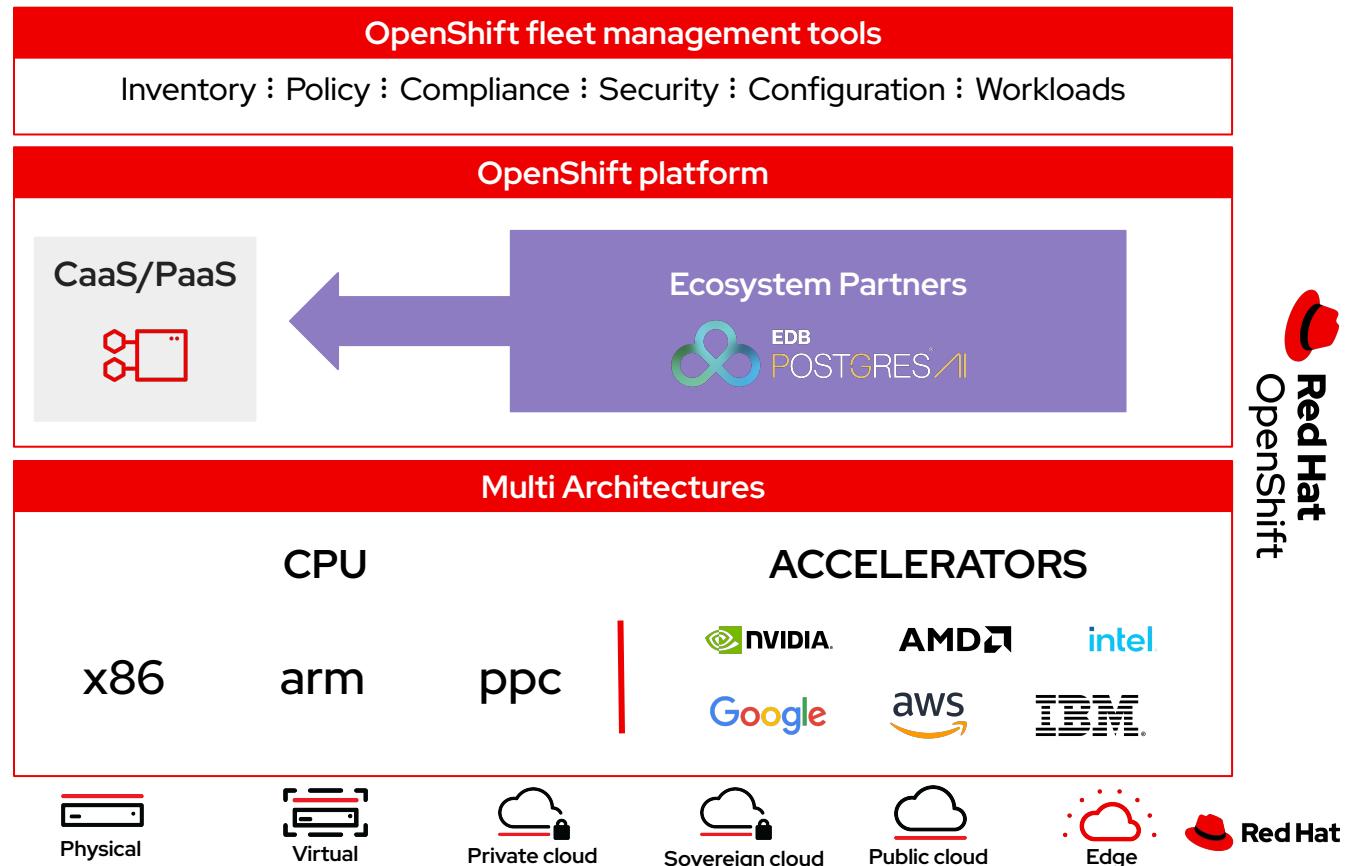
Red Hat's open source solutions work  
on bare metal and the public cloud ...

... and everything in between.

# A multi-faceted platform for strategic goals



**RED HAT OPENSHIFT**  
Container platform market  
share leader



# Red Hat OpenShift Application Platform

Trusted, comprehensive and consistent across hybrid cloud



# OpenShift is also **content**!

## Application runtimes

Included into OpenShift subscription (software collection)	PHP	Python	Java	NodeJS	Perl	Ruby	.NET Core	
	MySQL	PostgreSQL	MongoDB	Redis				
	Apache HTTP Server	Nginx	RH-SSO	JBoss Web Server	Tomcat			
	Spring Boot	Thorntail	Vert.x	JBoss EAP	JBoss AMQ Streams	JBoss AMQ	JBoss Fuse	
Available through additional bundles	3SCALE API mgmt	JBoss DAM	JBoss PAM	JBoss Data Virt	JBoss Data Grid	RH Mobile	JBoss Fuse Online	

... virtually any available container image !

## Marketplace

The screenshot shows the Red Hat OpenShift Software Catalog interface. On the left, there's a sidebar with navigation links like Home, Favorites, Ecosystem, Workloads, Virtualization, Migration, Serverless, Networking, Storage, Builds, Observe, Compute, User Management, and Administration. The main area is titled "Software Catalog" and contains a search bar and a list of items. Each item has a thumbnail, a name, and a brief description. A red box highlights the number "673 items" in the top right corner of the catalog area.

Category	Item Name	Description
All Items	AI/Machine Learning Application Runtime	Provided by Red Hat
All Items	Big Data Cloud Provider	Database
All Items	Cy/CD Database	Developer Tools
All Items	Cloud Provider Database	Drivers and plugins
All Items	Community .NET Application	Integration & Delivery
All Items	Community Helm Charts	Languages
All Items	Community [DEPRECATED] Quandrant Operator	Logging & Tracing
All Items	Community 3scale API Management	Middleware
All Items	Community 3scale Operator	Modernization & Migration
All Items	Community 3scale Follow	Monitoring
All Items	Community 3scale Operator	Networking
All Items	Community 3scale Operator	Observability
All Items	Community Advanced Cluster Security for Kubernetes	OpenShift Optional
All Items	Community Advanced Cluster Management for Kubernetes	Security
All Items	Community Advanced Cluster Security (RHACS) sponsor	Storage
All Items	Community Advanced Automation Platform	
All Items	Community An example Apache HTTP Server (httpd) application that serves static content. For more...	
All Items	Community Apache HTTP Server	
All Items	Community Helm Charts	

+600 items



# A large ecosystem

**AI / ML**

- GIGASPACE
- CognitiveScale
- PROPHETSTOR
- H<sub>2</sub>O
- SELDON
- UBIX
- PerceptiLabs

**Monitoring & Logging**

- BROADCOM
- ZABBIX
- INSTANA
- Red Hat
- turbonomic
- APPDYNAMICS part of Cisco
- sysdig
- dynatrace
- IBM
- sematext
- New Relic
- splunk>
- DATADOG

**Security**

- Twistlock
- CYBER ARMOR
- TREMOLO SECURITY
- aqua
- NeuVector
- JETSTACK
- portshift
- Zettaset
- IBM
- anchore
- tutin Orca
- SYNOPSIS

**Networking**

- CONTAINOUS
- TIGERA
- CITRIX®

**Customer Code**

{|}

**DevOps Tools**

- Kong
- ESCALATE Solution-Soft
- JFrog
- IBM VACNA
- sonatype
- Red Hat
- here
- OpsMx
- KubeMQ
- SPOT

**Application Runtimes**

- GTSoftware
- IBM
- Joget
- Red Hat
- Lightbend

**Databases & Big Data**

- memSQL
- hazelcast
- Starburst
- EDB
- redislabs HOME OF REDIS
- crunchydata
- ArangoDB
- NUODB
- PERCONA
- PingCAP
- Couchbase
- NoEQUAL
- mongoDB
- Cockroach LABS

**Storage**

- TRILIO
- INFINIDAT
- Red Hat
- portworx
- Hewlett Packard Enterprise
- ROBIN Storage
- kasten by Veeam
- NUTANIX STORAGEOS
- IBM

# aas platform for users

The screenshot shows the Red Hat OpenShift web interface with the title "aas platform for users". The top navigation bar includes the Red Hat logo, "Red Hat OpenShift", a dropdown for "local-cluster", and user information "dmartini@redhat.com". The left sidebar lists developer tools: "Developer", "+Add", "Topology", "Observe", "Search", "Functions", "Builds", "Pipelines", "Environments", "Helm", "Project", "ConfigMaps", and "Secrets". The main content area is titled "Project: dev-l2" and displays the "Developer Catalog". The catalog is organized into several sections:

- Developer Catalog**
  - All services**: Browse the catalog to discover, deploy and connect to services.
  - Database**: Browse the catalog to discover database services to add to your application.
  - Operator Backed**: Browse the catalog to discover and deploy operator managed services.
  - Helm Chart**: Browse the catalog to discover and install Helm Charts.
  - Virtual Machines**: Create a Virtual Machine from a template.
  - Samples**: Create an application from a code sample.
- Eventing**
  - Event Type**: Event Types available that can be consumed with Serverless Functions or regular Deployments.
  - Event Source**: Create an Event source to register interest in a class of events from a particular system.
  - Event Sink**: Create an Event sink to receive incoming events from a particular source.
  - Broker**: Create a Broker to define an event mesh for collecting a pool of events and route those events based on attributes, through triggers.
  - Channel**: Create a Knative Channel to create an event forwarding and persistence layer with in-memory and reliable implementations.
- OpenShift Self Managed Services**
  - Red Hat OpenShift AI**
- fn Serverless function**
  - Import from Git**: Create and deploy stateless, Serverless functions.
  - Samples**: Build from pre-built function sample code.
- Git Repository**
  - Import from Git**: Import code from your Git repository to be built and deployed.
- Container images**
  - Deploy an existing Image from an Image registry or Image stream tag.
- Sharing**
  - Project access allows you to add or remove a user's access to the project.
- From Local Machine**
  - Import YAML**: Create resources from their YAML or JSON definitions.
  - Helm Chart repositories**: Add a Helm Chart Repository to extend the Developer Catalog.
  - Upload JAR file**: Upload a JAR file from your local desktop to OpenShift.
- Pipelines**
  - Create a Tekton Pipeline to automate delivery of your application.

# Why Red Hat OpenShift for EDB: Operator Certification

The screenshot shows the Red Hat Ecosystem Catalog interface. At the top, there's a navigation bar with links for Solutions, Products, Artifacts, and Partners, along with a search bar labeled "Search Ecosystem Catalog". Below the navigation, a breadcrumb trail shows the user has navigated from Home to Software, then All software results, and finally Containerized applications. The main content area features a card for "EDB Postgres for Kubernetes", which is "Certified". The card includes a logo for EDB POSTGRES, a brief description: "PostgreSQL Operator for mission critical databases in Openshift Container Platform", and tabs for Overview, Resources, Certifications (which is selected), Deploy & use, and FAQs. Below this, a section titled "Certifications" contains a link to "Learn about Red Hat Certification and Partner Validation". Another section titled "Certified components" shows a list of container images, with the first item being "EDB Postgres for Kubernetes (formerly Cloud Native PostgreSQL Operator) Container Images" by "EnterpriseDB". This item includes a thumbnail, a green "Certified" badge, a description, and metadata like "Published 10 days ago". A search bar and pagination controls are also visible at the bottom of this section.

EDB Postgres for Kubernetes is a certified Level 5 Operator for Red Hat OpenShift

- ▶ This is designed to streamline Day 2 operations of PostgreSQL databases
- ▶ Enhanced Database Management
- ▶ Supports point-in-time recovery (PITR)
- ▶ Ensures robust data protection and recovery options
- ▶ Integration with business continuity solutions such as Red Hat OpenShift API for Data Protection (OADP) and Veeam Kasten, Trilio, Portworx Backup, IBM Fusion, and others

# Proven Solution at Scale



## EDB Postgres on OpenShift use cases

- Cloud-Native Database Deployment
- Database as a Service (DBaaS)
- High Availability and Disaster Recovery (HA & DR)
- DevOps and Continuous Integration/Continuous Deployment (CI/CD)
- Microservices and Application Modernization
- Move from VMWare to OpenShift
- Data Security and Compliance (using TDE and Advanced Security provided by EPAS)
- Hybrid and Multi-Cloud Deployments
- Multi-Tenant Applications (isolation)
- AI/ML Sovereign Platform

# Euro Information

## Company profile

Euro-Information is the fintech company of the Crédit Mutuel group. Euro-Information manages the IT systems of 16 federations of Crédit Mutuel as well as those of CIC and of all the financial, insurance, property, consumer credit, private banking, financing, telephony and technological subsidiaries.

## Problem

- Fast database deployment
- Adopt a supported and secure Open Source platform
- Onprem DBaaS
- Align to in-house RDBMS standardization

## Solution

- Use Postgres capabilities to build and maintain local applications
- Use Red Hat OpenShift platform to accelerate the provisioning of databases and applications

## Results

- Applications running with PostgreSQL databases in a centralized environment
- Massive reduction of TCO of database service operations



- Red Hat OpenShift
- EDB Postgres for Kubernetes
- PostgreSQL
- EPAS

- EDB considerably reduces IT costs associated with database maintenance.
- 280 cores: Enterprise Plan + Production Support

## Summary

Use Case

On prem DBaaS ([in Production](#))

Workload

Transactional

Application Name

All internal Postgres applications

EDB Tools of Interest

PostgreSQL and EDB Postgres for Kubernetes

# Airbus

## Company profile

Airbus SE is a European aerospace corporation. The company's primary business is the design and manufacturing of commercial aircraft but it also has separate defence and space and helicopter divisions.

## Problem

- Flexibility
- Cost reduction
- New managed service in OpenShift

## Solution

- EDB Postgres for Kubernetes with EPAS. Depending of the applications needs, EPAS and/or TDE will be activated

## Results

- POC done
- Decision is taken
- Number of cores not yet communicated



- Red Hat OpenShift
- EDB Postgres for Kubernetes
- EPAS
- TDE

- Improve database deployment speed
- Reduce DB support
- Cost reduction

## Summary

Use Case	On prem DBaaS ( <a href="#">Production</a> )
Workload	Transactional
Application Name	All VMWare PostgreSQL databases
EDB Tools of Interest	EDB Postgres Advanced Server with Oracle and TDE (optional)

# Banque de France

## Company profile

The Banque de France is France's central bank. A two-hundred-year-old institution, privately-owned when it was founded on January 18, 1800 under the Consulate by General Bonaparte, it became state-owned on January 1, 1946 when it was nationalized by General de Gaulle.

## Problem

- Fast database deployment
- Provide containerized Postgres DBaaS

## Solution

- Use OpenShift to provide this service with the operator
- Fast deployment and with Open Source database

## Results

- OpenShift based PostgreSQL cluster deployments expand the internal offering alongside traditional VM based database cluster deployments



- Red Hat OpenShift
- CloudNativePG
- PostgreSQL



- 100 cores
- Subscription plan:
  - Community360 plan + Production Support

## Summary

Use Case	OnPrem DBaaS ( <a href="#">In production</a> )
Workload	Transactional
Application Name	Multiple applications
EDB Tools of Interest	PostgreSQL, CloudNativePG



# La Poste

## Company profile

La Poste is a postal service company in France, operating in Metropolitan France, the five French overseas departments and regions and the overseas collectivity of Saint Pierre and Miquelon. Under bilateral agreements, La Poste also has responsibility for mail services in Monaco through La Poste Monaco and in Andorra alongside the Spanish company Correos.

## Problem

- Provide a database HA solution for Ansible Automation Platform (AAP)
- Database must be in HA and DR

## Solution

- Use EDB Postgres for Kubernetes to provide a HA and DR solution for PostgreSQL databases
- Deploy in 2 OpenShift clusters our operator

## Results

- La Poste developer can use their internal 'La Post Service Portal' to provision more than 64 backends.
- Reduce risk deploying EDB solutions.



- Red Hat OpenShift
- EDB Postgres for Kubernetes
- PostgreSQL

- EDB considerably reduces IT costs associated with database maintenance.
- 12 Cores: Standard Plan + Premium Support

## Summary

Use Case

On prem DBaaS with HA and DR  
[\(In Production\)](#)

Workload

Transactional

Application Name

Portail XaaS

EDB Tools of Interest

PostgreSQL and EDB Postgres for Kubernetes

# Thank you

Red Hat is the world's leading provider of enterprise open source software solutions.

Award-winning support, training, and consulting services make

Red Hat a trusted adviser to the Fortune 500.



[linkedin.com/company/red-hat](https://www.linkedin.com/company/red-hat)



[youtube.com/user/RedHatVideos](https://www.youtube.com/user/RedHatVideos)



[facebook.com/redhatinc](https://www.facebook.com/redhatinc)

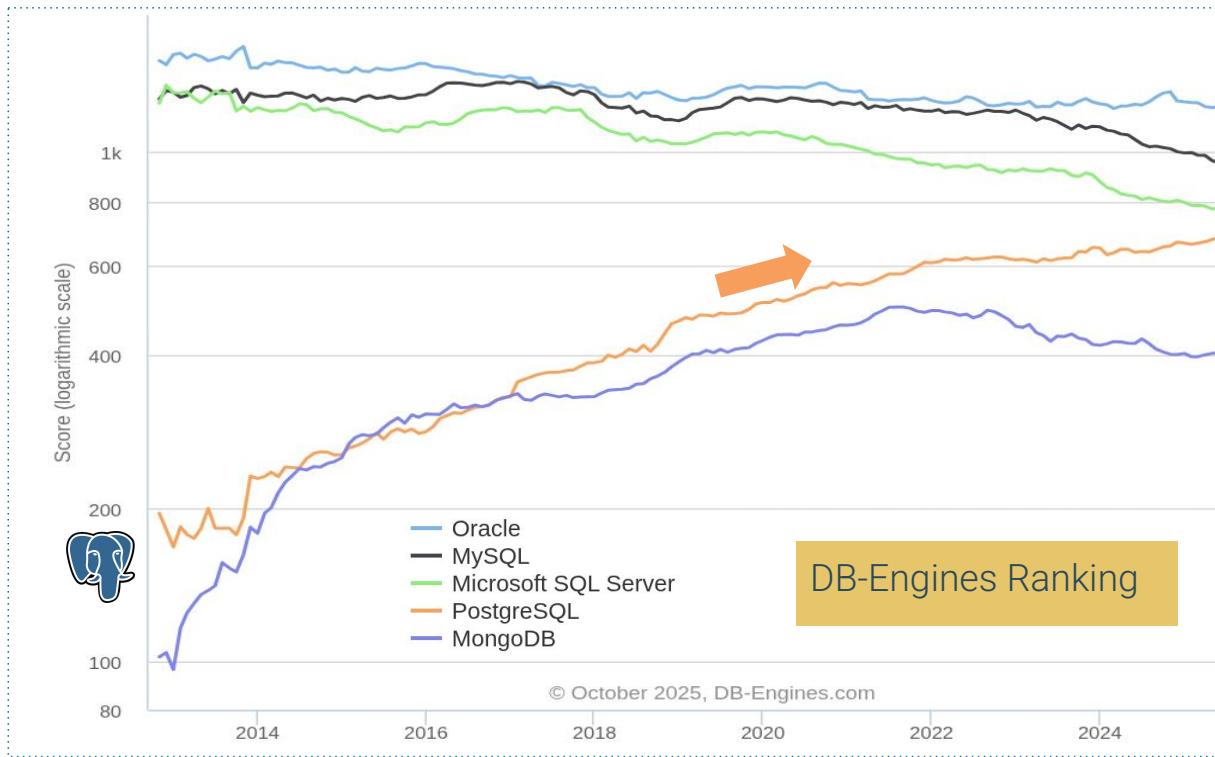


[twitter.com/RedHat](https://twitter.com/RedHat)

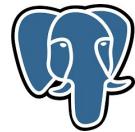
# Introduction to Postgres and EDB



# PostgreSQL is Booming !



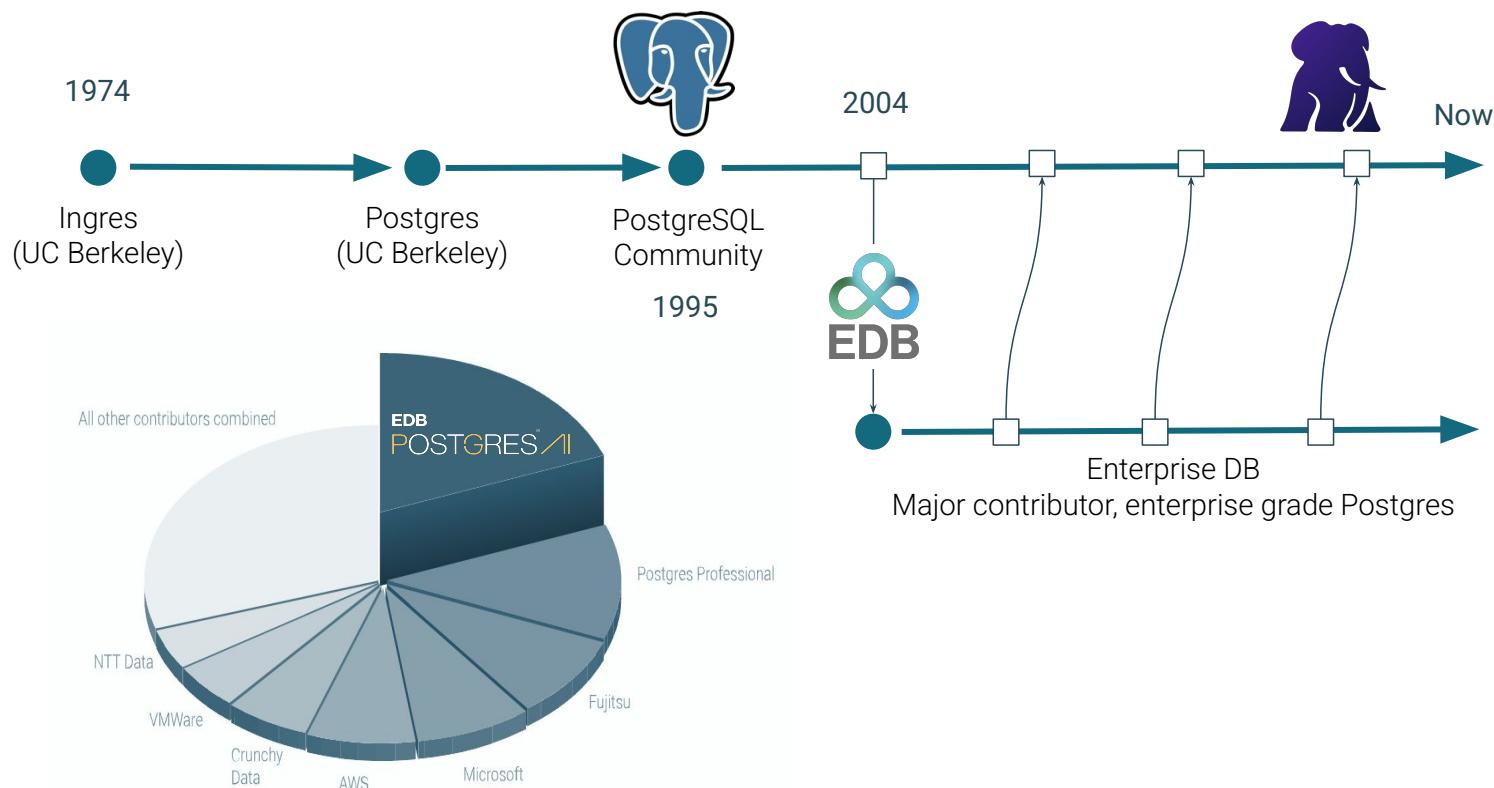
35%

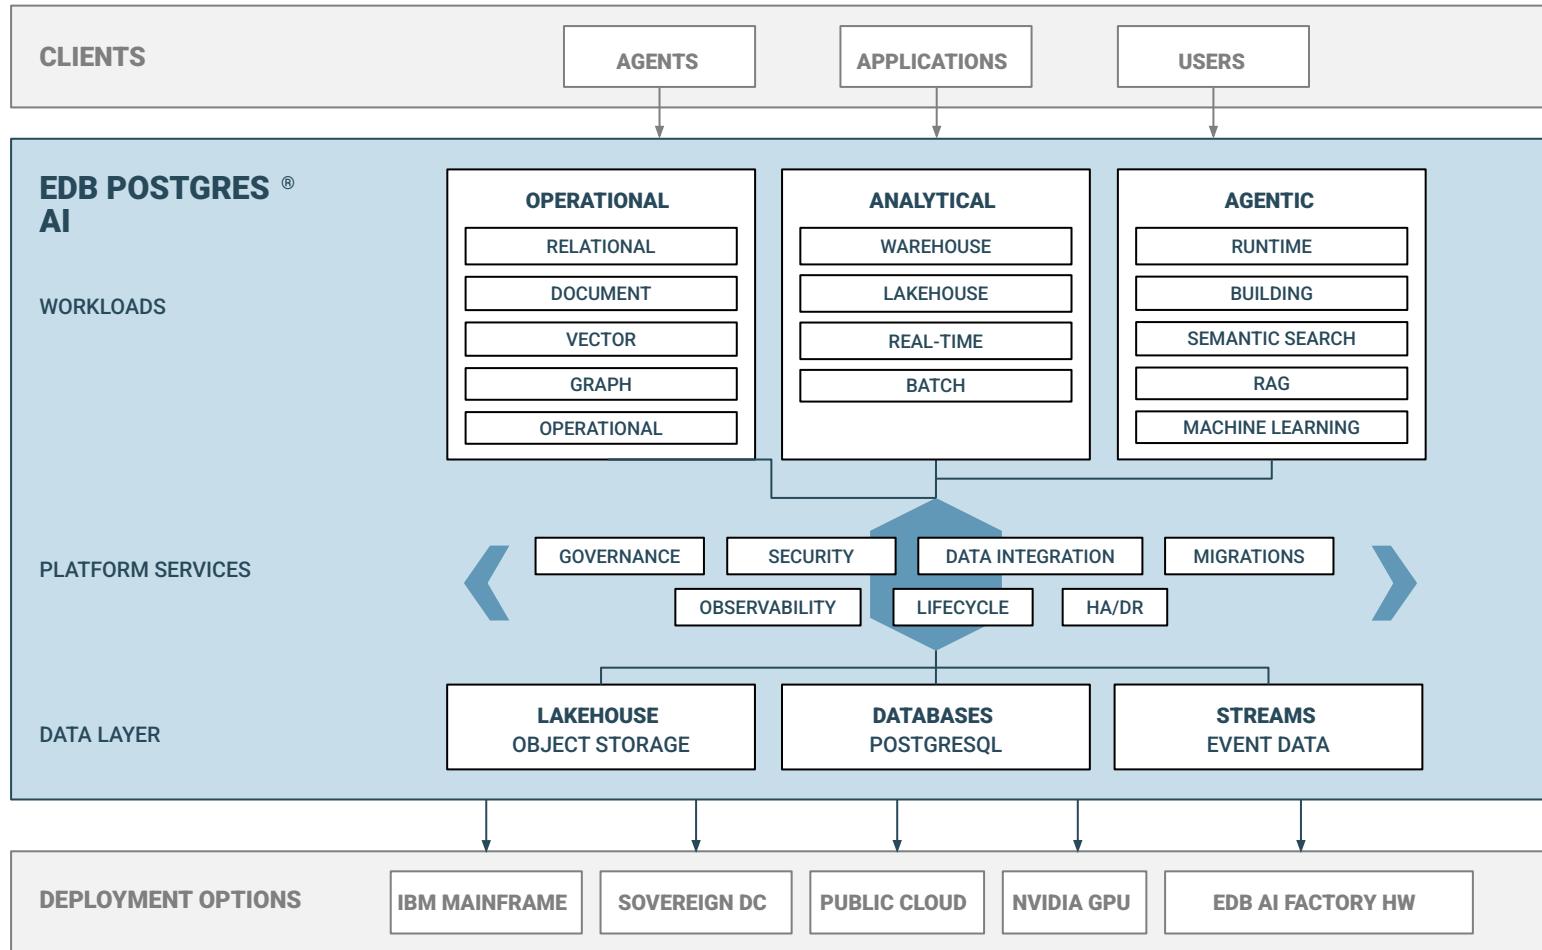


Thirty-five percent of Enterprises will consider Postgres for their next project



# EnterpriseDB Biggest PostgreSQL Contributor





# CNPG Operator: Reference Architecture and functionalities



# Kubernetes timeline

- 2014, June: Google open sources Kubernetes
- 2015, July: Version 1.0 is released
- 2015, July: Google and Linux Foundation start the CNCF
- 2016, November: The operator pattern is introduced in a blog post
- 2018, August: The Community takes the lead
- 2019, April: Version 1.14 introduces **Local Persistent Volumes**
- 2019, August: EDB team starts the Kubernetes initiative
- 2020, June: we publish this blog about benchmarking local PVs on bare metal
- 2020, June: Data on Kubernetes Community founded
- 2021, February: EDB Cloud Native Postgres (CNP) 1.0 released
- 2022, May: **EDB donates CNP** and open sources it under CloudNativePG
- 2025, January: CloudNativePG was recognized as an official **#CNCF** project



# A kubernetes operator for Postgres



Kubernetes adoption is rising and it is already the de facto **standard** orchestration tool



PostgreSQL clusters “management the kubernetes way” enables many cloud native usage patterns, e.g. spinning up, disposable clusters during tests, one cluster per microservice and one database per cluster



CNPG tries to encode years of experience managing PostgreSQL clusters into **an Operator which should automate all the known tasks** a user could be willing to do

Our PostgreSQL operator must simulate a part of the work of a DBA



## Autopilot

It automates the steps that a human operator would do to deploy and to manage a Postgres database inside Kubernetes, including automated failover.



## Security

A grayscale photograph of a security guard from behind. The guard is wearing a light-colored jacket with the word "SECURITY" printed in large, bold, white capital letters on the back. He is holding a dark walkie-talkie up to his ear with his right hand. The background is a plain, light-colored wall.

CloudNativePG is secured by default.





It doesn't rely on statefulsets and uses its own way to manage persistent volume claims where the PGDATA is stored.

## Data persistence



## Designed for Kubernetes

It's entirely declarative, and directly integrates with the Kubernetes API server to update the state of the cluster — for this reason, it does not require an external failover management tool.



# Features

Deployment	Administration	Backup & Recovery	Monitoring	Security	High Availability
Kubernetes operator	Single node	Backup	Prometheus	TDE	Switchover
Kubernetes plugin	Cluster (Multi node)	Recovery	Grafana dashboards	Certificates	Failover
EDB Postgres (EPAS)	PostgreSQL configuration	PITR	Postgres Enterprise Manager	Data redaction	Scale out / scale down
PostGIS	Pooling	Volume Snapshots	Logging	Password management	Minor / Major updates



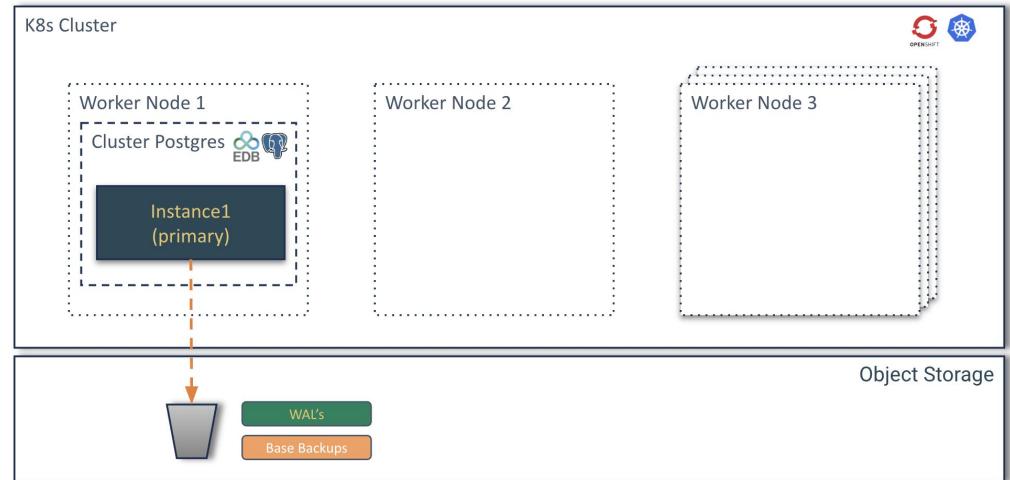
# Use Cases



# Use case 1 architecture

A single database is the simplest setup, involving one instance of a database server.

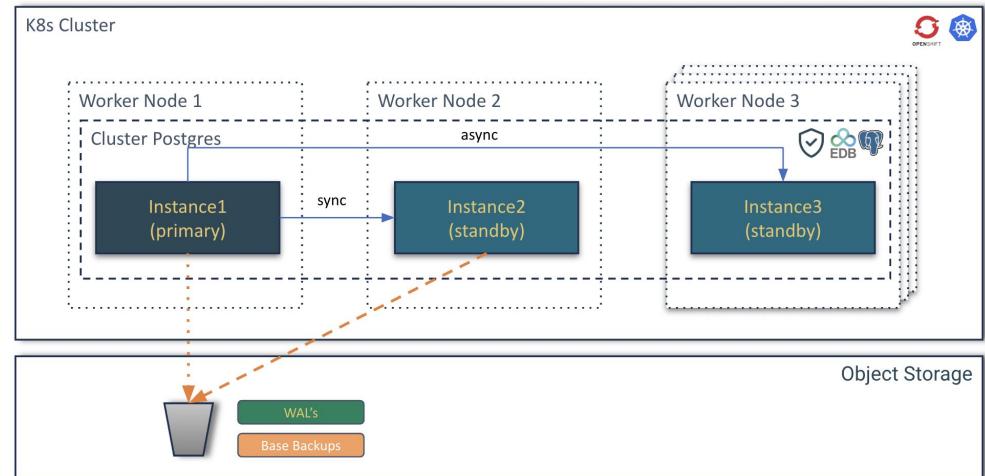
- Development and testing environments
- Small applications with low traffic
- Non-critical data analysis
- Applications with high tolerance for downtime
- Cost-sensitive projects



# Use case 2 architecture

An HA database setup aims to minimize downtime by having redundant components. If one component fails, another takes over automatically or with minimal intervention. This usually involves techniques like clustering, replication, or mirroring within the same data center or availability zone.

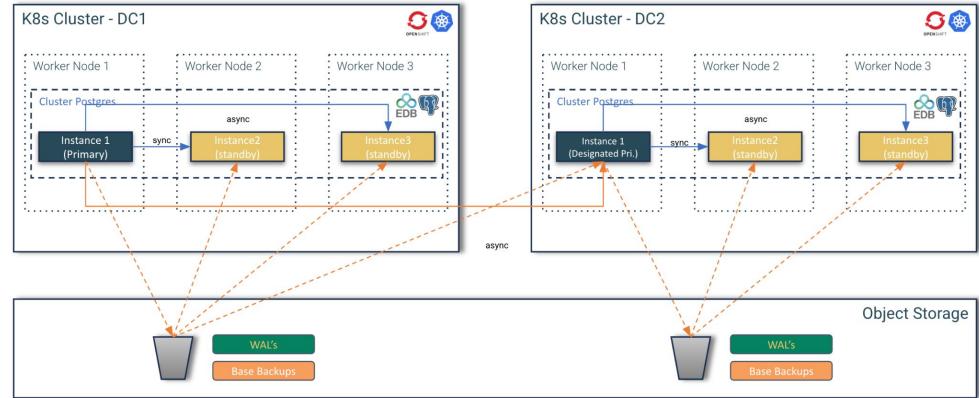
- Business critical Applications
- Applications with stringent SLAs
- Real-time systems
- Improving user experience
- Minimizing planned downtime



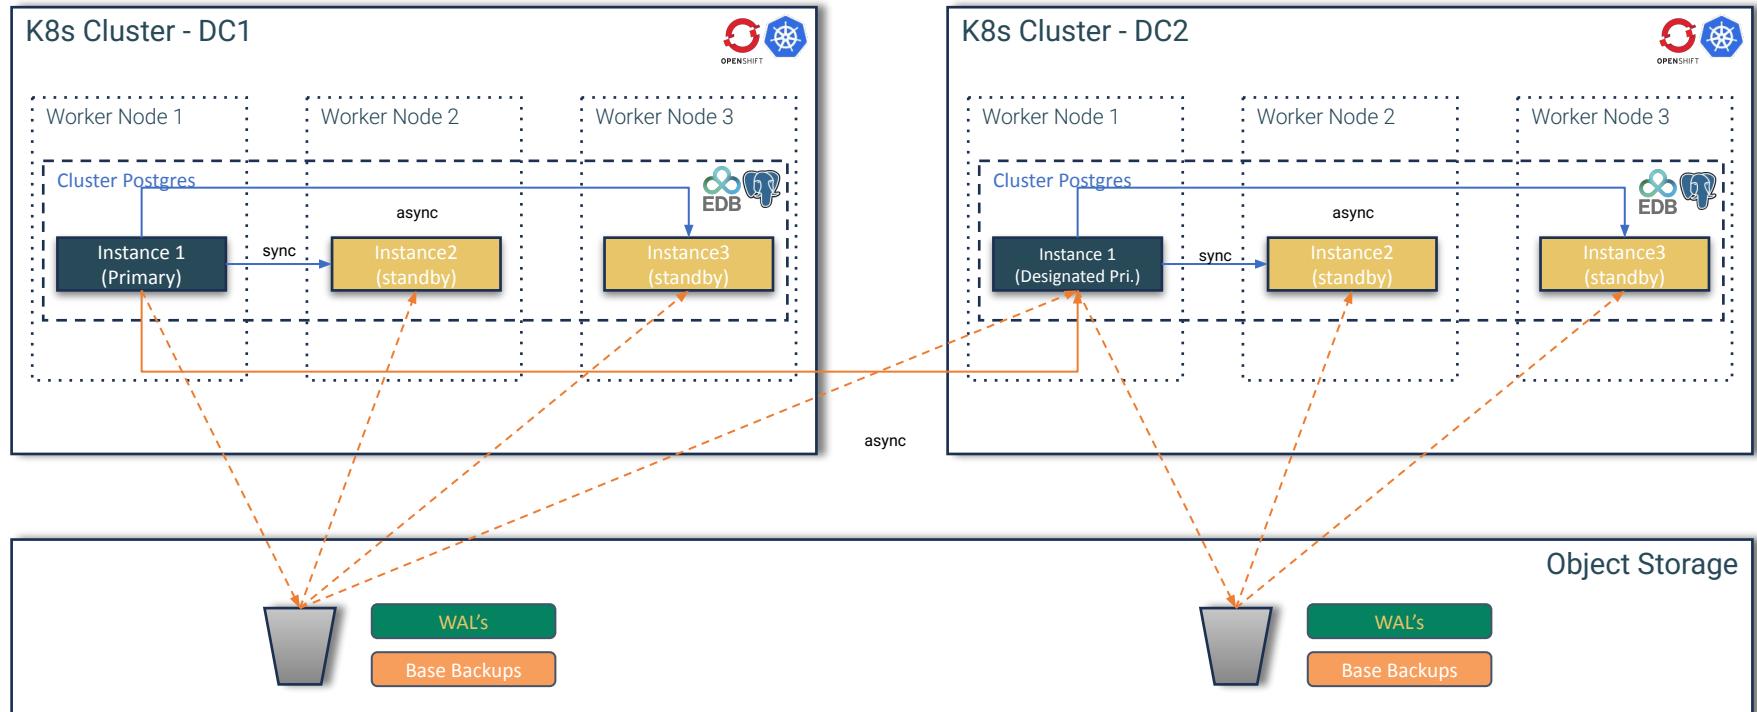
# Use case 3 architecture

A DR database setup focuses on protecting data and ensuring business continuity in the event of a large-scale disaster affecting an entire data center or region (e.g., natural disasters, power outages, cyberattacks). This typically involves replicating data to a geographically separate location.

- Regulatory compliance
- Protecting against catastrophic data loss
- Ensuring business continuity for mission-critical systems



# Use case 3 architecture



# Interactive session

# It's time to go hands-on!

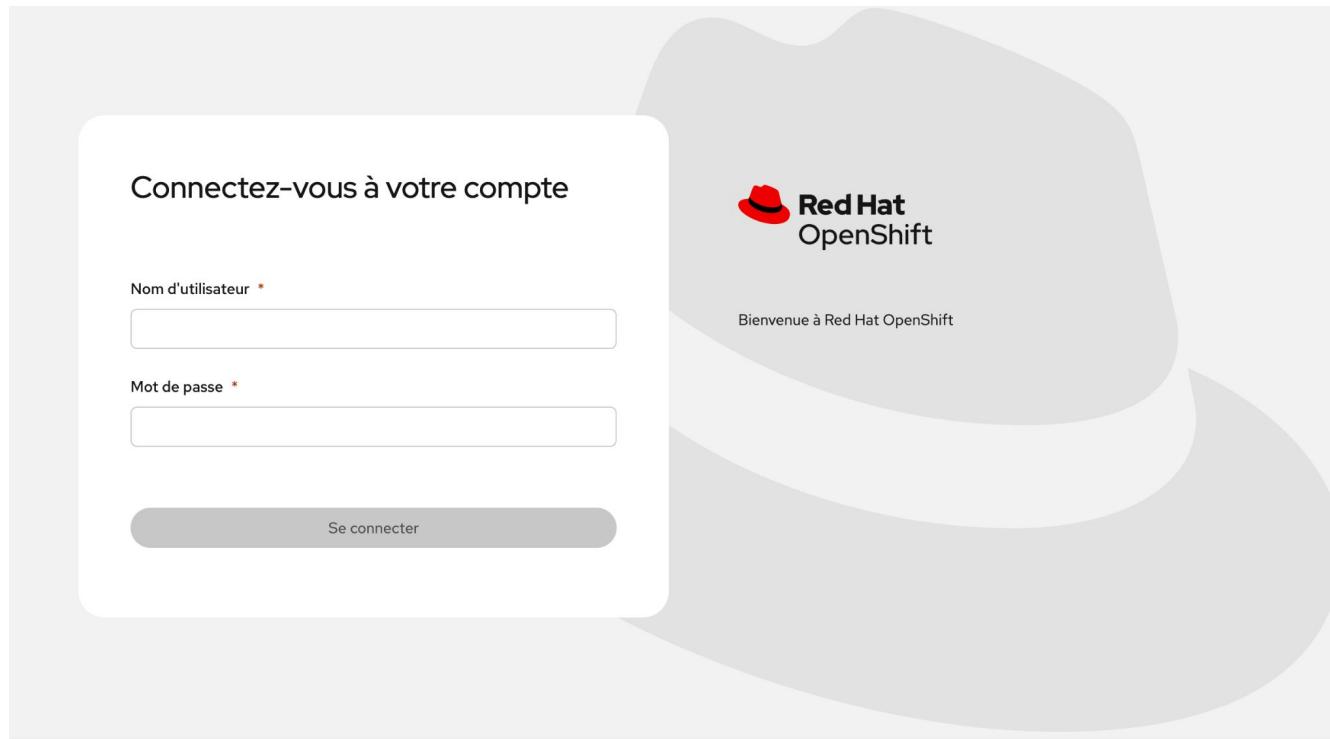


# Environment

<b>OpenShift Console</b>	<a href="https://console-openshift-console.apps.cluster-ldqnq.ldqnq.sandbox2860.opentlc.com">https://console-openshift-console.apps.cluster-ldqnq.ldqnq.sandbox2860.opentlc.com</a>
<b>OpenShift User name</b>	user1..user30
<b>OpenShift Password</b>	edbredhat
<b>Minio UI</b>	<a href="https://minio-ui-minio.apps.cluster-ldqnq.ldqnq.sandbox2860.opentlc.com/">https://minio-ui-minio.apps.cluster-ldqnq.ldqnq.sandbox2860.opentlc.com/</a>
<b>Minio internal service</b>	<a href="http://minio-service.minio.svc.cluster.local:9000">http://minio-service.minio.svc.cluster.local:9000</a>
<b>Minio API</b>	<a href="https://minio-api-minio.apps.cluster-ldqnq.ldqnq.sandbox2860.opentlc.com/">https://minio-api-minio.apps.cluster-ldqnq.ldqnq.sandbox2860.opentlc.com/</a>
<b>Minio User</b>	minio
<b>Minio password</b>	edb-workshop



# Red Hat OpenShift Login



# Red Hat OpenShift environment

The screenshot shows the Red Hat OpenShift web interface. The top navigation bar includes the Red Hat logo, a menu icon, a bell icon with the number 6, a plus sign, a question mark, and the user 'kube:admin'. A blue banner at the top states: 'You are logged in as a temporary administrative user. Update the [cluster OAuth configuration](#) to allow others to log in.' Below this, the 'Project: openshift-image-registry' is selected. The main content area is titled 'Installed Operators' with a star icon. It explains that operators are represented by ClusterServiceVersions and provides links to 'Understanding Operators documentation' and 'Operator SDK'. A search bar allows filtering by 'Name' or 'Search by name...'. A table lists installed operators, showing columns for Name, Managed Namespaces, Status, and Provided APIs. One operator, 'EDB Postgres for Kubernetes', is listed with details: version 1.28.0 provided by EDB, managed by namespace 'openshift-image-registry', and status 'Succeeded Up to date'. The 'Provided APIs' column includes links for Backups, Cluster Image Catalog, Cluster, Postgres Database, and View 6 more...

Name	Managed Namespaces	Status	Provided APIs
<a href="#">EDB Postgres for Kubernetes</a> 1.28.0 provided by EDB	<a href="#">NS openshift-image-registry</a> The operator is running in openshift-operators but is managing this namespace	Succeeded Up to date	<a href="#">Backups</a> <a href="#">Cluster Image Catalog</a> <a href="#">Cluster</a> <a href="#">Postgres Database</a> <a href="#">View 6 more...</a>



# Red Hat OpenShift operator

The screenshot shows the Red Hat OpenShift web console interface. The left sidebar navigation includes Home, Favorites, Ecosystem, Software Catalog, Installed Operators (selected), Helm, Workloads, Networking, Storage, Builds, Observe, Compute, User Management, and Administration. The main content area displays the 'Installed Operators' section for the 'openshift-image-registry' project. A specific operator, 'EDB Postgres for Kubernetes' (version 1.28.0 provided by EDB), is selected. The 'Operator details' tab is active. The page lists the following provided APIs:

Provider	Created at	Links
EDB	Dec 17, 2025, 10:38 AM	<a href="https://www.enterprisedb.com/products/postgresql-on-kubernetes-ha-clusters-k8s-containers-available">EDB Postgres for Kubernetes</a> <a href="https://www.enterprisedb.com/docs/postgres_for_kubernetes/latest/">Documentation</a>
Maintainers		Jonathan Gonzalez V. <a href="mailto:jonathan.gonzalez@enterprisedb.com">jonathan.gonzalez@enterprisedb.com</a> Jonathan Battato <a href="mailto:jonathan.battato@enterprisedb.com">jonathan.battato@enterprisedb.com</a> Niccolo Fei <a href="mailto:niccolo.fei@enterprisedb.com">niccolo.fei@enterprisedb.com</a> Gabriele Bartolini <a href="mailto:gabriele.bartolini@enterprisedb.com">gabriele.bartolini@enterprisedb.com</a>

**Provided APIs**

API Type	Description	Create Instance
Backups	PostgreSQL backup (physical base backup)	<a href="#">Create instance</a>
Cluster Image Catalog	A cluster-wide catalog of PostgreSQL operand images	<a href="#">Create instance</a>
Cluster	PostgreSQL cluster (primary/standby architecture)	<a href="#">Create instance</a>
Postgres Database	Declarative creation and management of a database on a Cluster	<a href="#">Create instance</a>
Failover Quorum	FailoverQuorum contains the information about the current failover quorum status of a PG cluster	<a href="#">Create instance</a>
Image Catalog	A catalog of PostgreSQL operand images	<a href="#">Create instance</a>
Pooler	Pooler for a Postgres Cluster (with PgBouncer)	<a href="#">Create instance</a>
Postgres Publication	Declarative creation and management of a Logical Replication Publication in a PostgreSQL Cluster	<a href="#">Create instance</a>
Scheduled Backups	Backup scheduler for a given Postgres cluster	<a href="#">Create instance</a>



# Use case

# The environment



# Features shown during the demo

- Kubernetes plugin install
- Check the CloudNativePG operator status
- Postgres cluster install
- Insert data in the cluster
- Failover
- Backup
- Recovery
- Scale out/down
- Fencing
- Hibernation
- Monitoring
- Rolling updates (minor and major)

Deployment

High Availability

Administration

Monitoring

Backup and Recovery

Last CloudNativePG tested version is 1.25



This demo is in



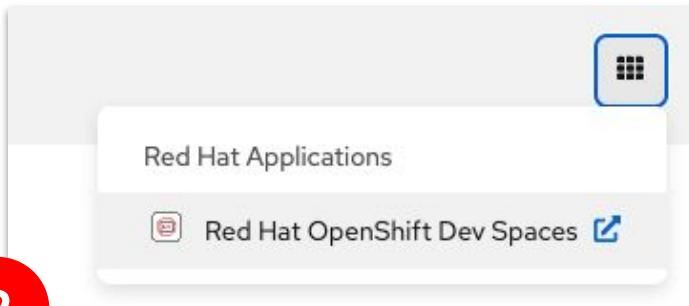
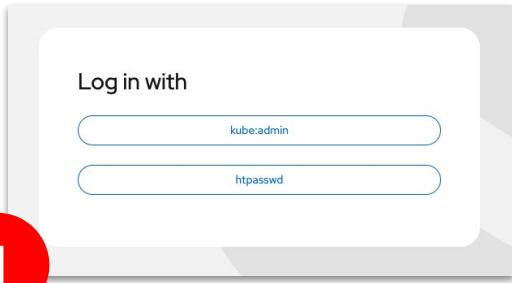
<https://github.com/raphael-chir/edb-postgres-for-kubernetes-in-openshift>



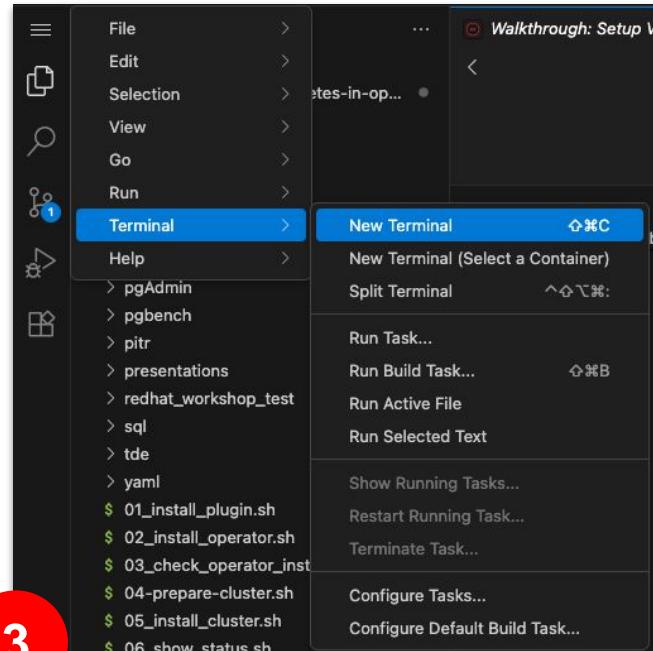
# Hands-on



# OpenShift and DevSpaces access



<https://github.com/raphael-chir/edb-postgres-for-kubernetes-in-openshift>



# Plugging and OpenShift login

## Scripts

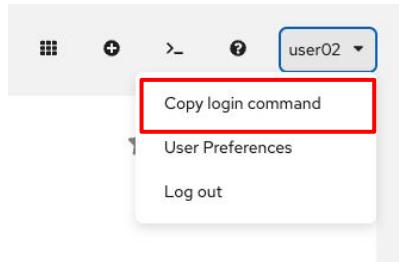
```
01_install_plugin.sh
```

### Install plugin

CloudNativePG provides a plugin for kubectl to manage a Postgres cluster in Kubernetes. You can manage status, promote, certificates, restart, reload, maintenance, report, logs, destroy, hibernation, backups, ...

Doc [link](#)

### OpenShift Login and namespace creation,



## Code

```
# Install plugin  
  
curl -sSfL \  
  https://github.com/EnterpriseDB/kubectl-cnp/raw/main/install.sh  
| sh -s -- -b .  
  
# Login to the OpenShift Cluster  
  
oc login ...  
  
# Create your namespace  
  
oc new-project userXX
```

# Checking operator installation

## Scripts

```
03_check_operator_installed.sh
```

### Checking operator installation

In Kubernetes, the operator is by default installed in the `postgresql-operator-system` namespace as a Kubernetes Deployment. The name of this deployment depends on the installation method. When installed through the manifest or the cnp plugin, by default, it is called `postgresql-operator-controller-manager`. When installed via Helm, by default, the deployment name is derived from the helm release name, appended with the suffix `-edb-postgres-for-kubernetes` (e.g., `<name>-edb-postgres-for-kubernetes`).

Doc [link](#)

## Code

```
# Check operator installation

oc get deploy -A | grep postgres
NAMESPACE: openshift-operators
NAME: postgresql-operator-controller-manager
READY: 1/1
UP-TO-DATE: 1
AVAILABLE: 1
AGE: 5d21h
```



# Install Postgres cluster

## Scripts

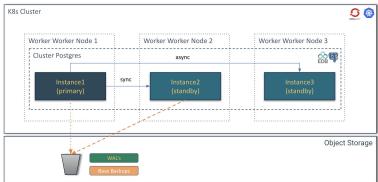
```
04-prepare-cluster.sh  
05_install_cluster.sh  
06_show_status.sh  
07_insert_data.sh
```

## Prepare Postgres cluster template

Setup credentials (AWS or Minio) and prepare context for your users

## Deploy Postgres cluster in HA

Different types of architectures can be deployed. In our demo, 3 instances are deployed with HA available.



## Insert Data

Insert data in test table

Doc [link](#)

## Code

```
export cluster_name="cluster-userXX"  
export TMP=$TMPDIR  
  
# Get template  
envsubst < templates/${cluster_name}-template.yaml >  
$TMP/${cluster_name}.yaml  
  
# Deploy Postgres cluster  
${kubectl_cmd} apply -f $TMP/${cluster_name}.yaml  
-----  
ApiVersion: postgresql.k8s.enterprisedb.io/v1  
kind: Cluster  
metadata:  
  name: cluster-example  
spec:  
  instances: 3  
  imageName: "docker.enterprisedb.com/k8s/edb-postgres-advanced:17.6"  
  enableSuperuserAccess: true  
  
replicationSlots:  
  highAvailability:  
    enabled: true
```



# Monitoring

## Show status

```
Cluster Summary
Name: default/cluster-sergio-1
System ID: 7566239654312464401
PostgreSQL Image: docker.enterprisedb.com/k8s/edb-postgres-advanced:17.6
Primary instance: cluster-sergio-1
Primary promotion time: 2025-10-28 11:52:15 +0000 UTC (4m30s)
Status: Cluster in healthy state
Instances: 3
Ready instances: 3
Size: 127M
Current Write LSN: 0/6032300 (Timeline: 1 - WAL File: 0000001000000000000006)

Continuous Backup status (Barman Cloud Plugin)
ObjectStore / Server name: object-store/cluster-sergio
First Point of Recoverability: -
Last Successful Backup: -
Last Failed Backup: -
Working WAL archiving: OK
WALs waiting to be archived: 0
Last Archived WAL: 00000010000000000005.0000028.backup @ 2025-10-28T11:54:11.908195Z
Last Failed WAL: -

Streaming Replication status
Replication Slots Enabled
Name Sent LSN Write LSN Flush LSN Replay LSN Write Lag Flush Lag Replay Lag State Sync State Sync Priority Replication Slot
---- -----
cluster-sergio-2 0/9000000 0/9000000 0/9000000 0/9000000 00:00:00 00:00:00 00:00:00 streaming quorum 1 active
cluster-sergio-3 0/9000000 0/9000000 0/9000000 0/9000000 00:00:00 00:00:00 00:00:00 streaming quorum 1 active

Instances status
Name Current LSN Replication role Status QoS Manager Version Node
---- -----
cluster-sergio-1 0/9000000 Primary OK Guaranteed 1.28.0 ocp4-multiarch-worker1-ppc64
cluster-sergio-2 0/9000000 Standby (sync) OK Guaranteed 1.28.0 ocp4-multiarch-w5gg4-master-2
cluster-sergio-3 0/9000000 Standby (sync) OK Guaranteed 1.28.0 ocp4-multiarch-w5gg4-master-1
```



# Promote and Failover

## Scripts

08\_promote.sh

### Promote

The meaning of this command is to promote a pod in the cluster to primary, so you can start with maintenance work or test a switch-over situation in your cluster:

Doc [link](#)

## Code

```
# Promote
. ./config.sh
${kubectl_cnp} promote ${cluster_name} ${cluster-name}-2
```



# Minor upgrade

## Scripts

09\_upgrade.sh

### Minor upgrade

To upgrade to a newer minor version, simply update the PostgreSQL container image reference in your cluster definition, either directly or via image catalogs. CloudNativePG will trigger a rolling update of the cluster, replacing each instance one by one, starting with the replicas. Once all replicas have been updated, it will perform either a switchover or a restart of the primary to complete the process.

Doc [link](#)

## Code

```
# Minor upgrade

# From
apiVersion: postgresql.k8s.enterprisedb.io/v1
kind: Cluster
metadata:
  name: cluster-sergio1
spec:
  instances: 3
  imageName:
"docker.enterprisedb.com/k8s/edb-postgres-advanced:17.6"
...

# To
apiVersion: postgresql.k8s.enterprisedb.io/v1
kind: Cluster
metadata:
  name: cluster-sergio1
spec:
  instances: 3
  imageName:
"docker.enterprisedb.com/k8s/edb-postgres-advanced:17.7"
...
```



# Backups

## Scripts

```
10_backup_cluster.sh  
11_backup_describe.sh
```

## Backup

Backups are used using Barman cloud with the new Barman Plugin. In case the plugin is not compatible (like in Red Hat OpenShift), we will continue using Barman Cloud (without the plugin).

PostgreSQL natively provides first class backup and recovery capabilities based on file system level (physical) copy. These have been successfully used for more than 15 years in mission critical production databases, helping organizations all over the world achieve their disaster recovery goals with Postgres.

Doc [link](#)

## Code

```
cat $TMP/backup.yaml  
  
apiVersion: postgresql.k8s.enterprisedb.io/v1  
kind: Backup  
metadata:  
  name: cluster-example-backup-test  
spec:  
  cluster:  
    name: cluster-example  
  
${{kubectl_cmd}} apply -f ./yaml/backup.yaml
```



# Recovery

## Scripts

```
12_restore_cluster.sh
```

## Recovery

In PostgreSQL, recovery refers to the process of starting an instance from an existing physical backup. PostgreSQL's recovery system is robust and feature-rich, supporting Point-In-Time Recovery (PITR)—the ability to restore a cluster to any specific moment, from the earliest available backup to the latest archived WAL file.

Doc [link](#)

## Code

```
cat $TMP/restore.yaml

apiVersion: postgresql.k8s.enterprisedb.io/v1
kind: Cluster
metadata:
  name: cluster-restore
spec:
  instances: 1
  imageName: quay.io/enterprisedb/postgresql:16.4
  imagePullPolicy: IfNotPresent
  bootstrap:
    recovery:
      source: source
  plugins:
  - name: barman-cloud.cloudnative-pg.io
    isWALArchiver: true
    parameters:
      barmanObjectName: object-store
  externalClusters:
  - name: source
    plugin:
      name: barman-cloud.cloudnative-pg.io
      parameters:
        barmanObjectName: object-store
        serverName: cluster-example

${kubectl_cmd} apply -f $TMP/restore.yaml
```



# Promote and Failover

## Scripts

13\_failover.sh

## Automated Failover

In the case of unexpected errors on the primary for longer than the `.spec.failoverDelay` (by default 0 seconds), the cluster will go into failover mode. This may happen, for example, when:

- The primary pod has a disk failure
- The primary pod is deleted
- The postgres container on the primary has any kind of sustained failure

In the failover scenario, the primary cannot be assumed to be working properly.

Doc [link](#)

## Code

```
# Automatic Failover (simulate failover killing pods of primary)
${kubectl_cmd} delete pvc/${primary} \
    pvc/${primary}-wal \
    pod/${primary} \
    --force
```



# Scale out, scale down

## Scripts

```
14_scale_out.sh  
15_scale_down.sh
```

## Scale out and down

The operator allows you to scale up and down the number of instances in a PostgreSQL cluster. New replicas are started up from the primary server and participate in the cluster's HA infrastructure.

Doc [link](#)

## Code

```
# Scale out  
kubectl scale cluster cluster-example --replicas=4
```

```
# Scale down  
kubectl scale cluster cluster-example --replicas=2
```



# Fencing and Hibernation

## Scripts

```
30_fencing_on.sh  
31_fencing_off.sh
```

## Fencing

Fencing is the process of protecting the data in one, more, or even all instances of a PostgreSQL cluster when they appear to be malfunctioning. When an instance is fenced, the PostgreSQL server process is guaranteed to be shut down, while the pod is kept running. This ensures that, until the fence is lifted, data on the pod isn't modified by PostgreSQL and that you can investigate file system for debugging and troubleshooting purposes.

Doc [link](#)

## Code

```
# Fencing on  
${kubectl_cnp} fencing on ${cluster_name} ${replica}  
  
# Fencing off  
${kubectl_cnp} fencing off ${cluster_name} ${replica}
```



# Fencing and Hibernation

## Scripts

```
32_hibernation_on.sh  
33_hibernation_off.sh
```

## Hibernation

CloudNativePG supports hibernation of a running PostgreSQL cluster in a declarative manner, through the cnpg.io/hibernation annotation. Hibernation enables saving CPU power by removing the database pods while keeping the database PVCs. This feature simulates scaling to 0 instances.

Doc [link](#)

## Code

```
# Hibernation on  
${{kubectl_cmd}} annotate cluster ${cluster_name} \  
--overwrite k8s.enterprisedb.io/hibernation=on  
  
# Hibernation off  
${{kubectl_cmd}} annotate cluster ${cluster_name} \  
--overwrite k8s.enterprisedb.io/hibernation=off
```



# Major upgrade (copying data)

## Scripts

20\_upgrade\_major\_version.sh

### Major upgrade (with native logical replication)

Major PostgreSQL releases introduce changes to the internal data storage format, requiring a more structured upgrade process.

CloudNativePG supports three methods for performing major upgrades:

- Logical dump/restore – Blue/green deployment, offline.
- Native logical replication – Blue/green deployment, online.
- **Physical with pg\_upgrade** – In-place upgrade, offline (covered in the "Offline In-Place Major Upgrades" section below).

Doc [link](#)

## Code

```
apiVersion: postgresql.k8s.enterprisedb.io/v1
kind: Cluster
metadata:
  name: ${cluster_major_upgrade}
spec:
  instances: 1
  imageName: ${postgres_major_upgrade_image}

  storage:
    size: 2Gi

  bootstrap:
    initdb:
      import:
        type: microservice
        databases:
          - app
        source:
          externalCluster: ${cluster_name}

  externalClusters:
    - name: ${cluster_name}
      connectionParameters:
        # Use the correct IP or host name for the source database
        host: ${cluster_name}-rw
        user: postgres
        dbname: postgres
        password:
          name: ${cluster_name}-superuser
          key: password
```



# Major upgrade (copying data)

## Scripts

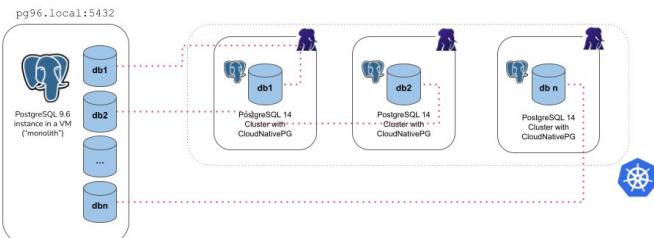
20\_upgrade\_major\_version.sh

### Major upgrade (with native logical replication)

- In this step we will migrate the app database from our existing cluster to the new cluster
- We will create the yaml file with the setting “**import**” in the bootstrap section
- The operator uses internally postgres tools pg\_dump and pg\_restore
- This method can be used to **migrate** another database or to run

#### out-of-the-place upgrade

- Possible settings:
  - Microservices
  - Monolith



## Code

```
apiVersion: postgresql.k8s.enterprisedb.io/v1
kind: Cluster
metadata:
  name: cluster-sergio-major
spec:
  instances: 1
  imageName: docker.enterprisedb.com/k8s/postgresql:18.1
  imagePullSecrets:
  - name: postgresql-operator-pull-secret
  enableSuperuserAccess: true

bootstrap:
  initdb:
    import:
      type: microservice
      databases:
      - postgres
      source:
        externalCluster: cluster-sergio

externalClusters:
- name: cluster-sergio
  connectionParameters:
    host: cluster-sergio-rw
    user: postgres
    dbname: postgres
    password:
      name: cluster-sergio-superuser
      key: password
  ...
```

# Nos prochains évènements

- **18 Mars** : Exclusive Virtual Wine Tasting
- **Mars/Avril** : Workshop/webinar en partenariat avec ACMI
- **23 Mars** : KubeCon + CloudNativeCon (Retex HSBC)
- **18 Juin** : Club Utilisateurs EDB (Retex Michelin)



**18 Mars**  
Exclusive Virtual  
Wine Tasting



**Mars/Avril**  
Workshop/webinar  
en partenariat  
avec ACMI



**23 Mars**  
KubeCon +  
CloudNativeCon  
(Retex HSBC)



**18 Juin**  
Club Utilisateurs  
EDB  
(Retex Michelin)

# Thank you

For any queries, contact:

[sergio.romera@enterprisedb.com](mailto:sergio.romera@enterprisedb.com)



# Backup Slides



# Architectures



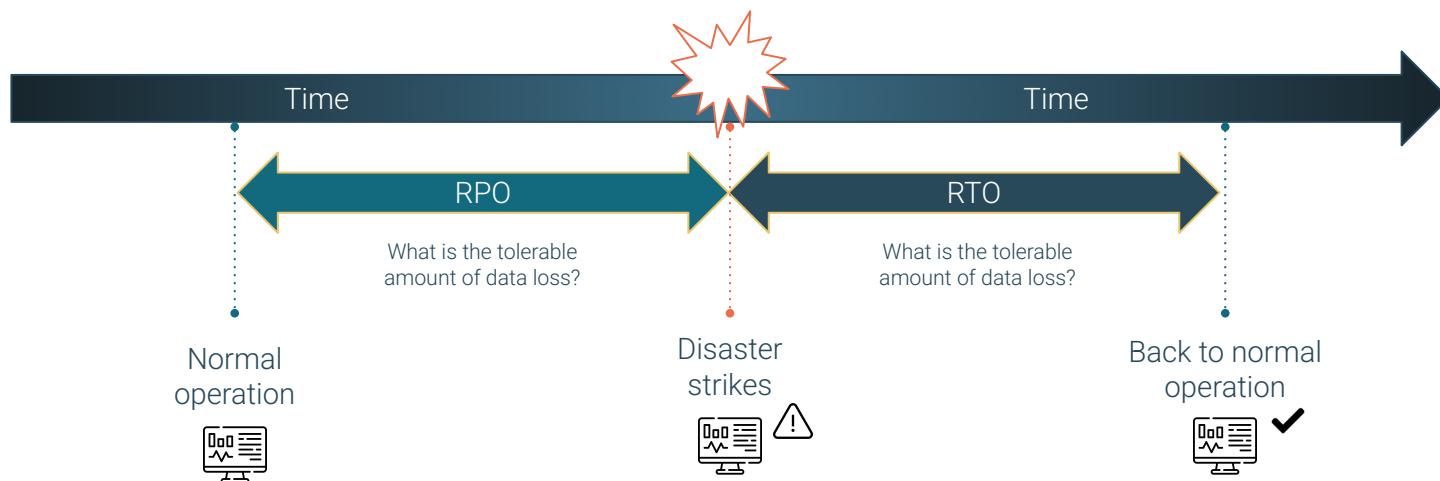
# When to choose Kubernetes over VMs?

- 01 |** Cloud Native Applications that already run in Kubernetes
- 02 |** Scalable, replicated databases
- 03 |** Applications requiring automated failover and self-healing
- 04 |** Teams skilled in Kubernetes who want a unified infrastructure



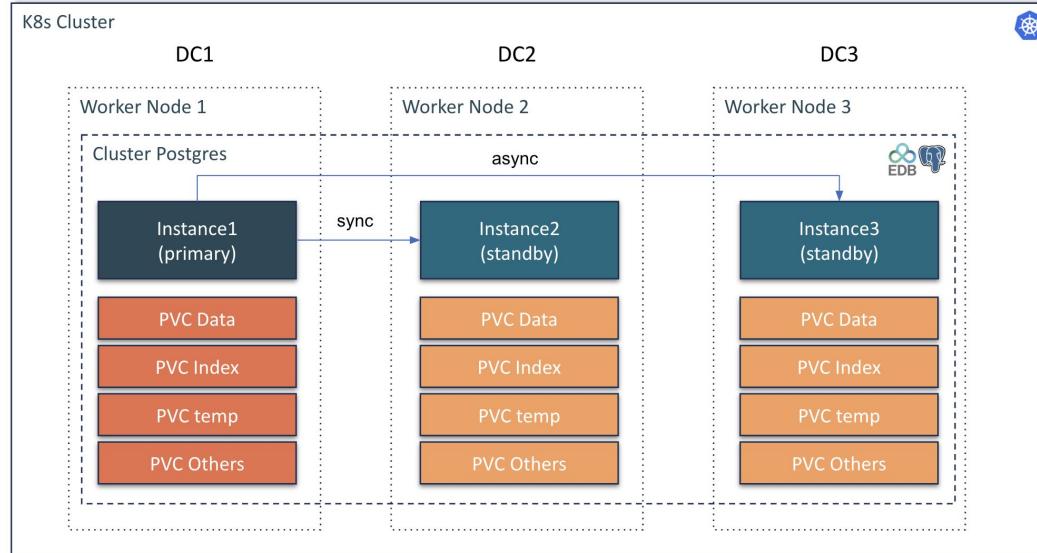
# Concepts

- Recovery Point Objective (**RPO**) and Recovery Time Objective (**RTO**) are key concepts in disaster recovery and business continuity planning, particularly related to data loss and system downtime.



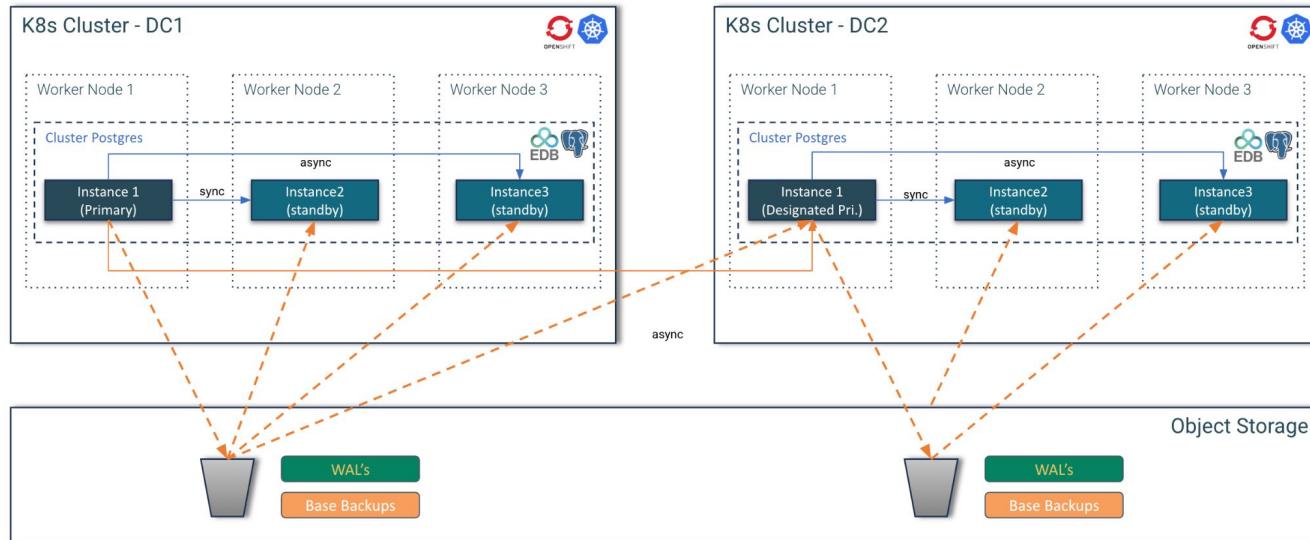
# Red Hat Recommendation

Red Hat recommend stretched clusters ONLY when latencies don't exceed 5 milliseconds (ms) round-trip time (RTT) between the nodes in different locations, with a maximum RTT of 10 ms.

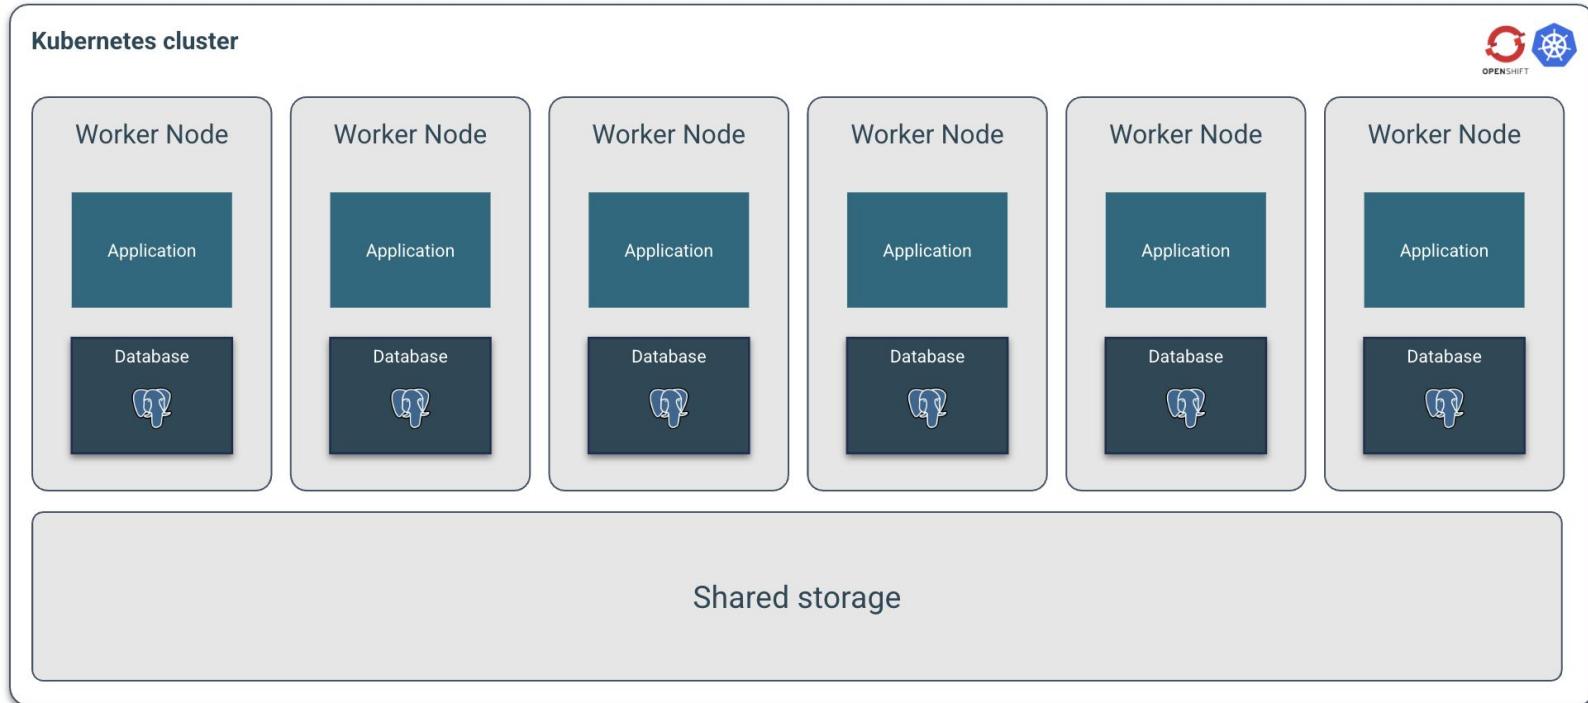


# Two separate single data center Kubernetes clusters

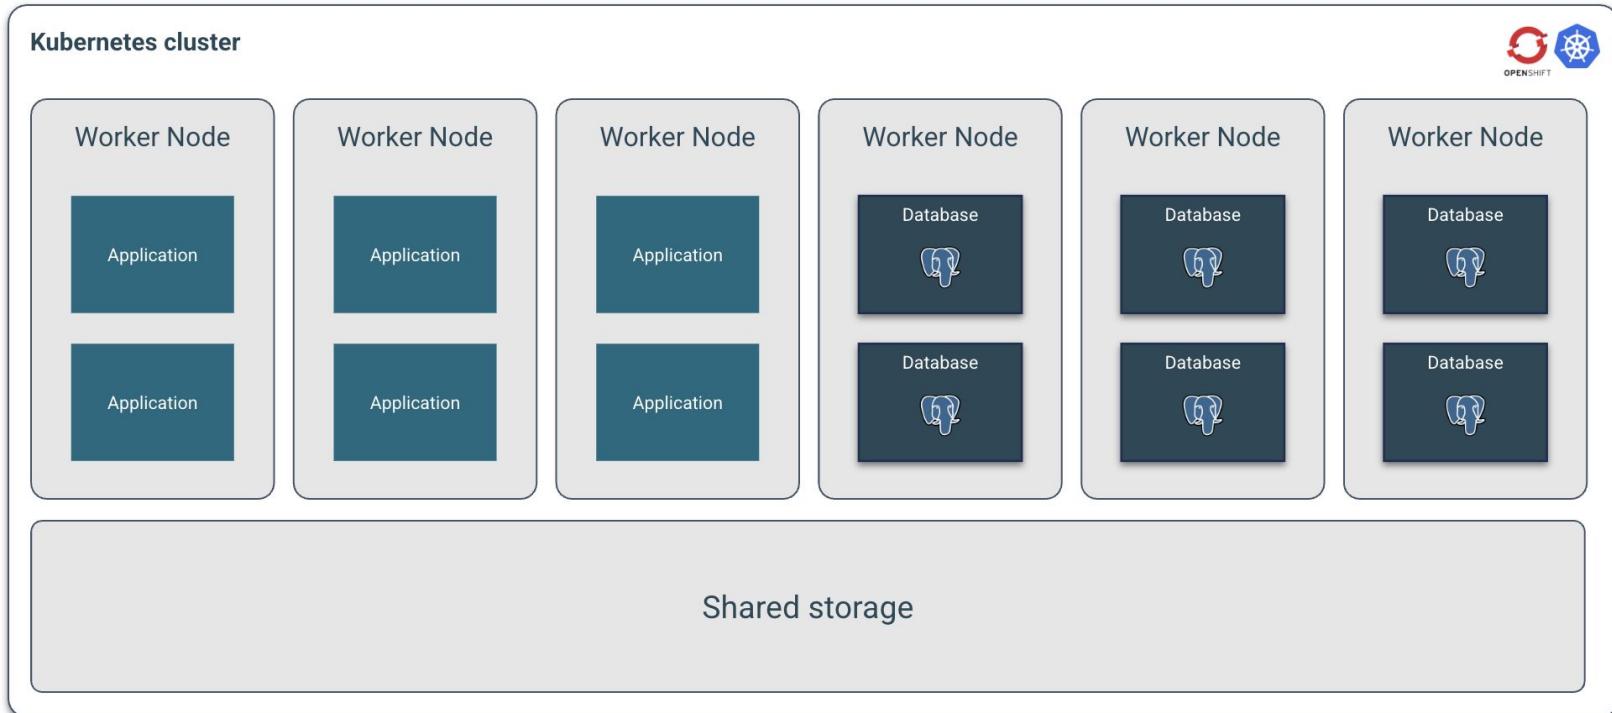
In case you cannot go beyond two data centers and you end up with two separate Kubernetes clusters, don't despair.



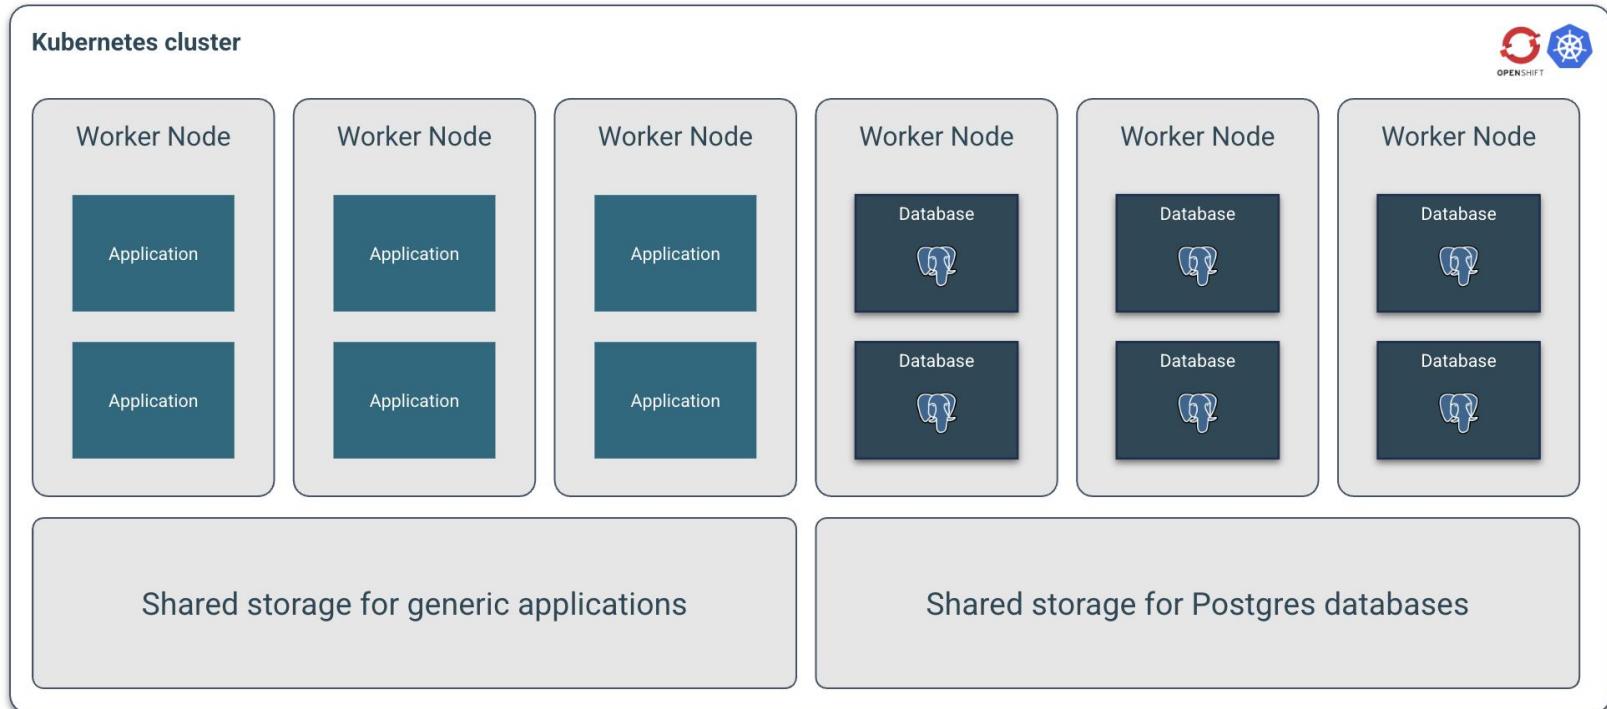
# Shared workload, shared storage



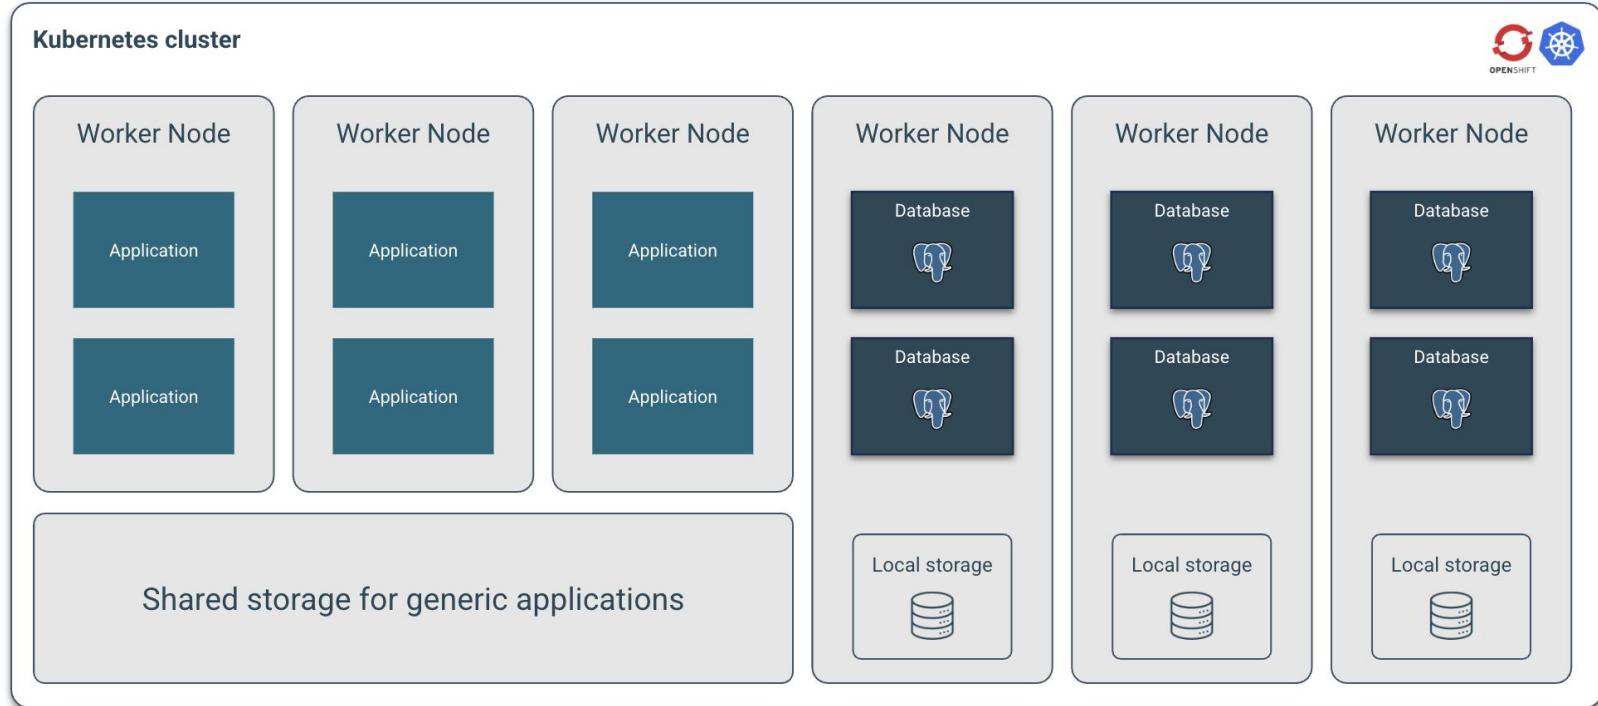
# Shared workload, shared storage



# Shared workload, shared storage



# Shared workloads, local storage



# Recommended architectures

<https://www.cncf.io/blog/2023/09/29/recommended-architectures-for-postgresql-in-kubernetes/>



## Recommended architectures for PostgreSQL in Kubernetes

By Gabriele Bartolini

September 29, 2023

Member post by Gabriele Bartolini, VP of Cloud Native at EDB

"You can run databases on Kubernetes because it's fundamentally the same as running a database on a VM", [tweeted Kelsey Hightower just a few months ago](#). Quite the opposite from what the former Google engineer and advocate said back in 2018 on Twitter: "Kubernetes supports stateful workloads; I don't."



Kelsey Hightower  
@kelseyhightower

You can run databases on Kubernetes because it's fundamentally the same as running a database on a VM. The biggest challenge is understanding that rubbing Kubernetes on Postgres won't turn it into Cloud SQL. ■

Truth is that I agree with him now as much as I agreed with him back then. At that time, the holistic offering of storage capabilities in Kubernetes was still immature (local persistent volumes would become GA only the year after), the operator pattern – which in the meantime has proven to be crucial for stateful applications like databases – was yet to become widely accepted, and the [Data on Kubernetes Community](#) was more than two years away (second half of 2020).

Nowadays, the situation is completely different. And I am sure that many people who've worked hard in the last few years to bring stateful workloads in Kubernetes agree with me that Kelsey's recent powerful words will contribute to reversing the public perception and facilitate our mission – provided we keep doing great.



# PostgreSQL most loved database (2023)

Source: [Stackoverflow](#)

This year, PostgreSQL took over the first place spot from MySQL. Professional Developers are more likely than those learning to code to use PostgreSQL (50%) and those learning are more likely to use MySQL (54%).

MongoDB is used by a similar percentage of both Professional Developers and those learning to code and it's the second most popular database for those learning to code (behind MySQL).



# Links:

## Openshift Console:

<https://console-openshift-console.apps.cluster-826x2.826x2.sandbox5464.opentlc.com/>

## Users:

name: user2..user32  
Password: edb-workshop

## Devspaces url:

<https://devspaces.apps.cluster-826x2.826x2.sandbox5464.opentlc.com>

## Short url to Devspaces:

<https://tinyurl.com/mrxs95pb>

## Minio:

UI: <https://minio-ui-minio.apps.cluster-826x2.826x2.sandbox5464.opentlc.com>  
API: <https://minio-api-minio.apps.cluster-826x2.826x2.sandbox5464.opentlc.com>

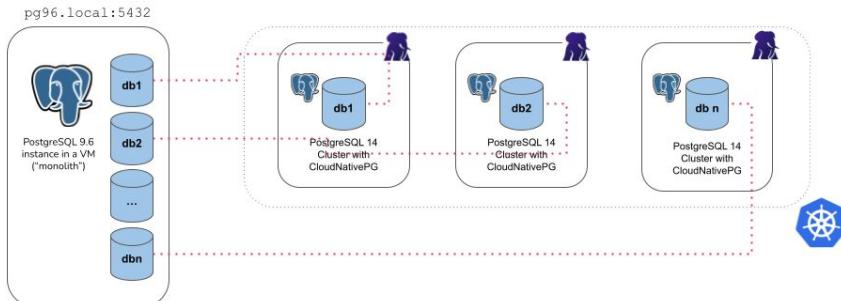
User: minio  
Password: edb-workshop



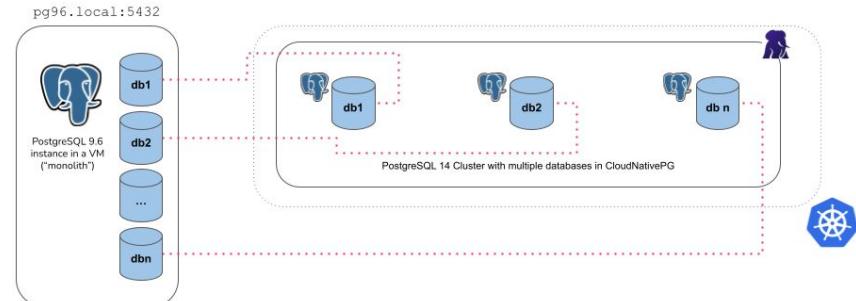
# Database Migration

- In this step we will migrate the app database from our existing cluster to the new cluster
- We will create the yaml file with the setting “**import**” in the bootstrap section
- The operator uses internally postgres tools pg\_dump and pg\_restore
- This method can be used to **migrate** another database or to run **out-of-the-place upgrade**
- Possible settings:

Microservices:



Monolith:



# Major upgrade (copying data)

## Scripts

```
24_major_upgrade_in_place.sh  
25_verify_major_upgrade_16_17.sh
```

### Major upgrade (offline In-Place Major Upgrades)

CloudNativePG performs an offline in-place major upgrade when a new operand container image with a higher PostgreSQL major version is declaratively requested for a cluster.

Major upgrades are only supported between images based on the same operating system distribution. For example, if your previous version uses a bullseye image, you cannot upgrade to a bookworm image.

Doc [link](#)

## Code

```
apiVersion: postgresql.k8s.enterprisedb.io/v1  
kind: Cluster  
metadata:  
  name: ${cluster_major_upgrade}  
spec:  
  instances: 1  
  imageName: ${postgres_major_upgrade_image}  
  
storage:  
  size: 2Gi  
  
bootstrap:  
  initdb:  
    import:  
      type: microservice  
      databases:  
        - app  
...  
...
```



# Disaster recovery

## Scripts

```
./dr/01-create-all.sh  
./dr/02-insert-eu.sh  
./dr/03-switchover-cluster.sh
```

## Disaster recovery

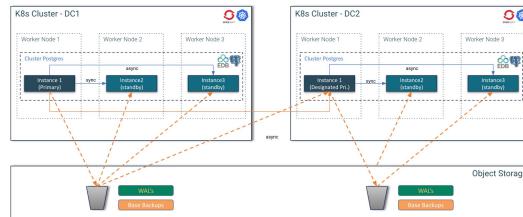
A replica cluster is a CloudNativePG Cluster resource designed to replicate data from another PostgreSQL instance, ideally also managed by CloudNativePG.

Typically, a replica cluster is deployed in a different Kubernetes cluster in another region. These clusters can be configured to perform cascading replication and can rely on object stores for data replication from the source, as detailed further down..

### Two clusters:

- pg-eu
- pg-us

Doc [link](#)



## Code

```
# EU  
kubectl delete objectstores.barmancloud.cnpg.io object-storage-eu  
envsubst < ../templates/dr/azure-object-storage-eu.yaml >  
.tmp/object-storage-eu.yaml  
kubectl apply -f ./tmp/object-storage-eu.yaml  
sleep 2  
rm -f ./tmp/pg-eu.yaml  
envsubst < ../templates/dr/pg-eu.yaml > ./tmp/pg-eu.yaml  
print_info "Creating pg-eu cluster...\n"  
kubectl apply -f ./tmp/pg-eu.yaml  
print_info "pg-eu cluster created.\n"  
  
# US  
kubectl delete objectstores.barmancloud.cnpg.io object-storage-us  
envsubst < ../templates/dr/azure-object-storage-us.yaml >  
.tmp/object-storage-us.yaml  
kubectl apply -f ./tmp/object-storage-us.yaml  
sleep 2  
rm -f ./tmp/pg-us.yaml  
envsubst < ../templates/dr/pg-us.yaml > ./tmp/pg-us.yaml  
kubectl wait --timeout=30m --for=condition=Ready cluster/pg-eu  
print_info "Creating pg-us cluster...\n"  
kubectl apply -f ./tmp/pg-us.yaml  
./backup_cluster.sh pg-eu  
print_info "pg-us cluster created.\n"
```

# The “cnp” plugin for kubectl

- The official CLI for CloudNativePG
  - Available also as RPM or Deb package
- Extends the ‘kubectl’ command:
  - Customize the installation of the operator
  - Status of a cluster
  - Perform a manual switchover (promote a standby) or a restart of a node
  - Issue TLS certificates for client authentication
  - Declare start and stop of a Kubernetes node maintenance
  - Destroy a cluster and all its PVC
  - Fence a cluster or a set of the instances
  - Hibernate a cluster
  - Generate jobs for benchmarking via pgbench and fio
  - Issue a new backup
  - Start pgadmin



Name: cluster-example  
Namespace: default  
System ID: 7100921006673293335  
PostgreSQL Image: ghcr.io/cloudnative-pg/postgresql:14.3  
Primary instance: cluster-example-2  
Status: Cluster in healthy state  
Instances: 3  
Ready instances: 3  
Current Write LSN: 0/C000060 (Timeline: 4 - WAL File: 00000004000000000000000C)

#### Certificates Status

Certificate Name	Expiration Date	Days Left Until Expiration
cluster-example-replication	2022-08-21 13:15:00 +0000 UTC	89.95
cluster-example-server	2022-08-21 13:15:00 +0000 UTC	89.95
cluster-example-ca	2022-08-21 13:15:00 +0000 UTC	89.95

#### Continuous Backup status

First Point of Recoverability: 2022-05-23T13:37:08Z  
Working WAL archiving: OK  
WALs waiting to be archived: 0  
Last Archived WAL: 00000004000000000000000B @ 2022-05-23T13:42:09.37537Z  
Last Failed WAL: -

#### Streaming Replication status

Name	Sent LSN	Write LSN	Flush LSN	Replay LSN	Write Lag	Flush Lag	Replay Lag	State	Sync State	Sync Priority
cluster-example-3	0/C000060	0/C000060	0/C000060	0/C000060	00:00:00	00:00:00	00:00:00	streaming	async	0
cluster-example-1	0/C000060	0/C000060	0/C000060	0/C000060	00:00:00	00:00:00	00:00:00	streaming	async	0

#### Instances status

Name	Database Size	Current LSN	Replication role	Status	QoS	Manager Version
cluster-example-3	33 MB	0/C000060	Standby (async)	OK	BestEffort	1.15.0
cluster-example-2	33 MB	0/C000060	Primary	OK	BestEffort	1.15.0
cluster-example-1	33 MB	0/C000060	Standby (async)	OK	BestEffort	1.15.0



# Install CNPG plugin

NOT NEEDED DURING WORKSHOP  
For illustrative purposes.

- In the web terminal run the script 01\_install\_plugin.sh:

```
./01_install_plugin.sh
```



# Call to action: Check CNPG plugin

- Call the help for the CNPG Plugin, run:

```
kubectl-cnp help
```



# Use case Operator installation



# Operator Installation demonstration

NOT NEEDED DURING WORKSHOP  
For illustrative purposes.

- Install the operator
- Check the installed CNP Operator in the console
- Discover the features of the Operator in the OpenShift environment
- Check the installed CNP Operator in the web terminal



# Install the CNPG Operator and check the it in the web terminal

NOT NEEDED DURING WORKSHOP  
For illustrative purposes.

- In the web terminal install the operator:

```
./02_install_operator.sh
```

**-> will require admin privs on Openshift**



## Call to action: Check the it in the web terminal

- In the web terminal check the installation of the operator:

```
./03_check_operator_installed.sh
```



# Call to action: Check the installed CNPG Operator in the Openshift console

- In the OpenShift console navigate to:
  - -> Operators
  - -> Installed Operators
  - -> Klick on the Operator installed:

The screenshot shows the Red Hat OpenShift web interface. On the left, there's a sidebar with navigation links: Home, Favorites, Operators (with a dropdown), OperatorHub, Installed Operators (which is selected and highlighted in white), Helm, Workloads, Networking, Storage, Builds, and Observe. The main content area has a header 'Project: All Projects' and a sub-header 'Installed Operators'. It says 'Installed Operators are represented by ClusterServiceVersions within this Namespace.' Below this is a search bar with 'Name' and 'Search by name...' fields. A table lists operators: 1. 'cert-manager Operator for Red Hat OpenShift' (Namespace: cert-manager-operator, Status: Succeeded, Up to date, Last updated: 14. Okt. 2025, 12:19, APIs: CertificateRequest, Certificate, Challenge, ClusterIssuer, View 4 more...). 2. 'EDB Postgres for Kubernetes' (Namespace: openshift-operators, Status: Succeeded, Up to date, Last updated: 15. Okt. 2025, 05:43, APIs: Backups, Cluster Image Catalog, Cluster, Postgres Database, View 6 more...). A green arrow points to the second row.

Name	Namespace	Managed Namespaces	Status	Last updated	Provided APIs
<a href="#">cert-manager Operator for Red Hat OpenShift</a>	<a href="#">NS cert-manager-operator</a>	<a href="#">NS cert-manager-operator</a>	<span>Succeeded</span> Up to date	14. Okt. 2025, 12:19	<a href="#">CertificateRequest</a> <a href="#">Certificate</a> <a href="#">Challenge</a> <a href="#">ClusterIssuer</a> <a href="#">View 4 more...</a>
<a href="#">EDB Postgres for Kubernetes</a>	<a href="#">NS openshift-operators</a>	All Namespaces	<span>Succeeded</span> Up to date	15. Okt. 2025, 05:43	<a href="#">Backups</a> <a href="#">Cluster Image Catalog</a> <a href="#">Cluster</a> <a href="#">Postgres Database</a> <a href="#">View 6 more...</a>

# Call to action: Discover the features of the Operator in the OpenShift environment

Project: openshift-operators ▾

[Installed Operators](#) ➤ Operator details



EDB Postgres for Kubernetes

1.27.0 provided by EDB

◀ Details YAML Subscription Events All instances Backups Cluster Image Catalog Cluster Postgres Database Failover Quorum

## Provided APIs

### B Backups

PostgreSQL backup (physical base backup)

[+ Create instance](#)

### CIC Cluster Image Catalog

A cluster-wide catalog of PostgreSQL operand images

[+ Create instance](#)

### C Cluster

PostgreSQL cluster (primary/standby architecture)

[+ Create instance](#)

### D Postgres Database

Declarative creation and management of a database on a Cluster

[+ Create instance](#)

### FQ Failover Quorum

FailoverQuorum contains the information about the current failover quorum status of a PG cluster

[+ Create instance](#)

### IC Image Catalog

A catalog of PostgreSQL operand images

[+ Create instance](#)

### P Pooler

### P Postgres Publication

### SB Scheduled Backups

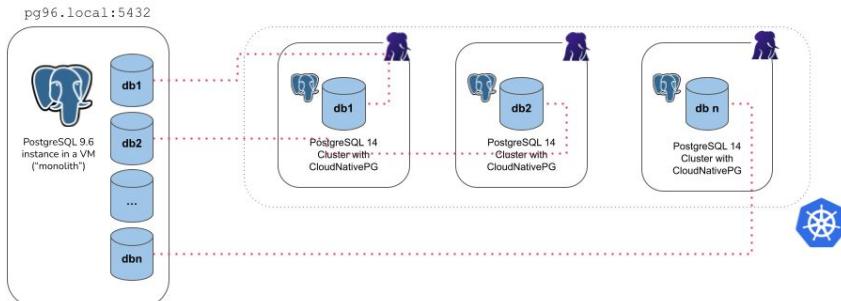
# Use case Database Migration / Major Version Upgrade



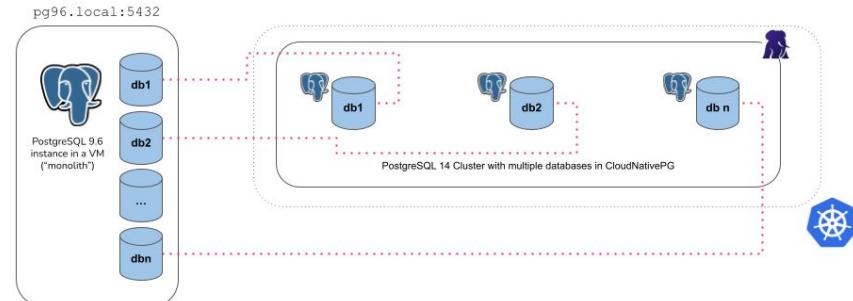
# Database Migration

- In this step we will migrate the app database from our existing cluster to the new cluster
- We will create the yaml file with the setting “**import**” in the bootstrap section
- The operator uses internally postgres tools pg\_dump and pg\_restore
- This method can be used to **migrate** another database or to run **out-of-the-place upgrade**
- Possible settings:

Microservices:



Monolith:

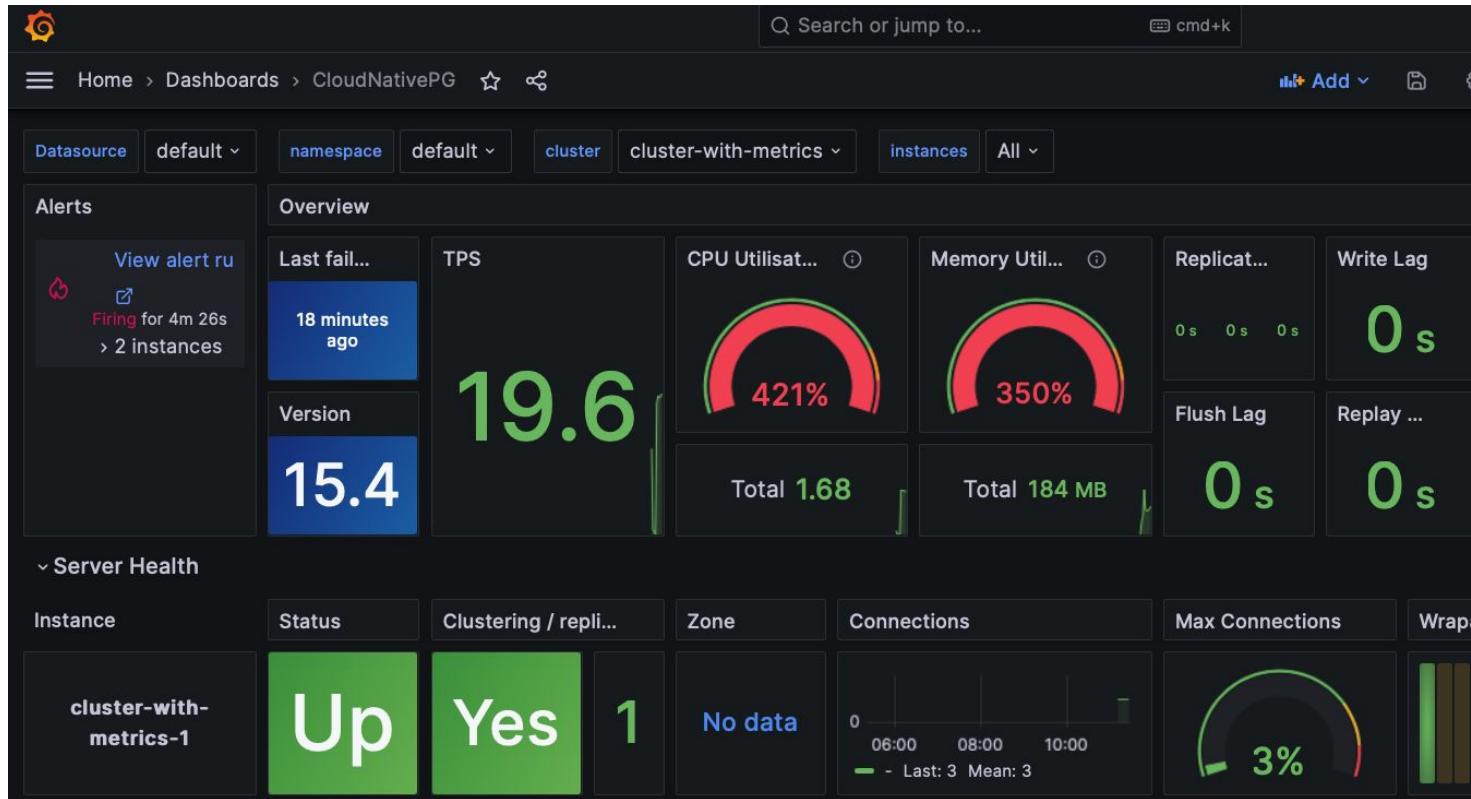


# Call to action: Run migration or out of the place upgrade

- In the web terminal 1:
  - Run the script:  
./20\_upgrade\_major\_version.sh
  - Check the cluster creation process:  
kubectl get pods -w
  - Check the table test in the cluster-user<X>-17, run the command:
    - Connect to the cluster:  
oc cnp psql cluster-user<X>-17
    - Connect to the database app:  
\c app
    - Run sql commands:  
select version();  
select count(\*) from test;



# Grafana Dashboard



# Agenda

Start	End	Session
12:30	13.30	Registration & Welcome Lunch
13:30	13:45	Campnocamp Vorstellung
13:45	14:00	Red Hat OpenShift & EDB Partnerschaft
14:00	14.15	Einführung in den Postgres-Marketplace
14:15	14:30	Coffee
14:30	15:00	CNPG Operator Reference Architecture
15:00	15:30	Functionalities for CNPG
15:30	17:15	Interactive session & demo
17:15	18:30	Drinks & Apero

