

POSTGRES ON OPENSIFT WORKSHOP



David Martini - OpenShift Specialist, Red Hat
Raphaël Chir - Senior Sales Engineer, EDB
Sidy-Mohamed Aidara- Account Executive, EDB

26th May, 2026



Agenda

Commencer	Fin	Session
12:30	13:30	Accueil Cocktail Lunch
13:30	13:45	Partenariat entre Red Hat OpenShift et EDB (Red Hat - David Martini)
13:45	14:00	Présentation du marché Postgres et d'EDB (EDB - Eric Pillon)
14:00	14:30	Architecture de référence de l'opérateur CNPG (EDB - Raphaël Chir)
14:30	16:30	Atelier pratique (EDB - Lucie Zeng & Raphaël Chir)
16:30	17:00	Pause
17:00	17:30	Conclusion et perspectives



Red Hat OpenShift with EDB

David Martini
Senior Solution Architect

The world's leading provider of open source enterprise IT solutions

More than
90%
of the
Fortune
500
use
Red Hat
products and
solutions¹

~22,000
employees

105+
offices

The first
\$3
billion
open
source
company
in the world²

Sources: 1-Red Hat client data and [Fortune 500 list](#), September 2024.


2 - [Red Hat SEC filings](#) prior to the acquisition by IBM. Note: Currency in U.S. dollars.

Open source

Open source is about more than developing software.
It's how we built our company.

And it's why we have been so successful.

From communities to enterprise



Linux contributions



-  Platinum members
-  Red Hat has been working on Linux since 1993
-  Linux is the foundation of all Red Hat products
-  The first Patent promise was issued by Red Hat in 2002

OpenInfra contributions




-  Historical member
-  First Red Hat OpenStack release in 2013
-  Most represented member on the board (3/21)
-  Red Hat & IBM = more than 3 times contributions as the second contributor

Cloud-native contributions



-  Co-founder in 2015
-  Platinum members
-  Second Kubernetes contributor after Google
-  First CNCF contributor in 2022, 2023, 2024, etc...

CNCF 2024 Global contribution



Red Hat is the **#1** company for contributions to **CNCF** in **2024**.

Red Hat portfolio

Red Hat's open source solutions work
on bare metal and the public cloud ...

... and everything in between.

A multi-faceted platform for strategic goals

OpenShift fleet management tools

Inventory : Policy : Compliance : Security : Configuration : Workloads

OpenShift platform

CaaS/PaaS



Ecosystem Partners



Multi Architectures

CPU

ACCELERATORS

x86

arm

ppc



Physical



Virtual



Private cloud



Sovereign cloud



Public cloud




Edge



Red Hat

Red Hat
OpenShift



RED HAT OPENSIFT
Container platform market
share leader

Red Hat OpenShift Application Platform

Trusted, comprehensive and consistent across hybrid cloud

Application Build



Developer Tools



Runtimes & languages



Serverless & functions



Cloud Services APIs



Migration tools

Application Delivery



CI/CD & GitOps



Container registry



Observability



Logging



Virtual machines

Application Security



Service mesh



Vulnerability scanning



Security policies



Cert & secret tooling



Auth and SSO services

OpenShift is also content !

Application runtimes

Included into OpenShift subscription (software collection)	PHP	Python	Java	NodeJS	Perl	Ruby	.NET Core	Third-party Language Runtimes
	MySQL	PostgreSQL	MongoDB	Redis	... virtually any available container image !			Third-party Databases
	Apache HTTP Server	Nginx	RH-SSO	JBoss Web Server				Third-party App Runtimes
Available through additional bundles	Spring Boot	Thorntail	Vert.x	JBoss EAP	JBoss AMQ Streams	JBoss AMQ	JBoss Fuse	Third-party Middleware
	3SCALE API mgmt	JBoss DAM	JBoss PAM	JBoss Data Virt	JBoss Data Grid	RH Mobile	JBoss Fuse Online	Third-party Middleware

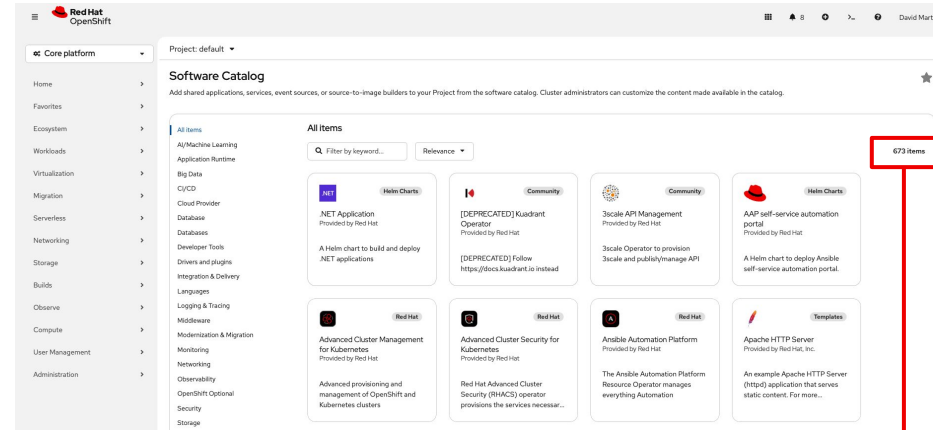
Health Index ⓘ



This image does not have any unapplied Critical or Important security updates.





The Container Health Index analysis is based on RPM packages signed and created by Red Hat, and does not grade other software that may be included in a container image.

Marketplace



+600 items

A large ecosystem



aaS platform for users

local-cluster

Developer

+Add

Topology

Observe

Search

Functions

Builds

Pipelines

Environments

Helm

Project

ConfigMaps

Secrets

Project: dev-l2

All services

Browse the catalog to discover, deploy and connect to services

Database

Browse the catalog to discover database services to add to your application

Operator Backed

Browse the catalog to discover and deploy operator managed services

Helm Chart

Browse the catalog to discover and install Helm Charts

Virtual Machines

Create a Virtual Machine from a template

Samples

Create an application from a code sample

Eventing

Event Type

Event Types available that can be consumed with Serverless Functions or regular Deployments

Event Source

Create an Event source to register interest in a class of events from a particular system

Event Sink

Create an Event sink to receive incoming events from a particular source

Broker

Create a Broker to define an event mesh for collecting a pool of events and route those events based on attributes, through triggers

Channel

Create a Knative Channel to create an event forwarding and persistence layer with in-memory and reliable implementations

Serverless function

Import from Git

Create and deploy stateless, Serverless functions

Samples

Build from pre-built function sample code

Sharing

Project access allows you to add or remove a user's access to the project

Helm Chart repositories

Add a Helm Chart Repository to extend the Developer Catalog

Pipelines

Create a Tekton Pipeline to automate delivery of your application

Git Repository

Import from Git

Import code from your Git repository to be built and deployed

Container images

Deploy an existing Image from an Image registry or Image stream tag

From Local Machine

Import YAML

Create resources from their YAML or JSON definitions

Upload JAR file

Upload a JAR file from your local desktop to OpenShift

OpenShift Self Managed Services

Red Hat OpenShift AI

Why Red Hat OpenShift for EDB: Operator Certification

The screenshot shows the Red Hat Ecosystem Catalog interface. At the top, there's a navigation bar with the Red Hat logo, 'Red Hat Ecosystem Catalog', and links for Solutions, Products, Artifacts, and Partners. A search bar is also present. Below the navigation bar, a breadcrumb trail reads: Home > Software > All software results > Containerized applications. The main content area features a large card for 'EDB Postgres for Kubernetes' with a 'Certified' badge. Below the card, there are tabs for Overview, Resources, Certifications (which is active), Deploy & use, and FAQs. Under the 'Certifications' tab, there's a section titled 'Certifications' with a link to 'Learn about Red Hat Certification and Partner Validation'. Below that, a section titled 'Certified components' includes a search bar, an 'Image type' dropdown, and a pagination indicator '1 - 2 of 2'. The main list of components shows 'EDB Postgres for Kubernetes (formerly Cloud Native PostgreSQL Operator) Container Images' with a green 'A' icon, the image name 'enterprisedb/cloud-native-postgresql', a description, and metadata including 'Container image', '1.26.0-rc2-ubi9', and 's390x'. To the right of the list, there's a sidebar with 'EnterpriseDB' and 'Published 10 days ago'.

Red Hat Ecosystem Catalog Solutions Products Artifacts Partners All Search Ecosystem Catalog

Home > Software > All software results > Containerized applications

EDB POSTGRES/11

EDB Postgres for Kubernetes

PostgreSQL Operator for mission critical databases in Openshift Container Platform

Certified

Overview Resources **Certifications** Deploy & use FAQs

Certifications

Learn about Red Hat Certification and Partner Validation

Certified components

Search Image type 1 - 2 of 2

EDB POSTGRES/11

EDB Postgres for Kubernetes (formerly Cloud Native PostgreSQL Operator) Container Images

enterprisedb/cloud-native-postgresql

Container images for EDB Postgres for Kubernetes (formerly Cloud Native PostgreSQL operator)

Container image 1.26.0-rc2-ubi9 s390x

EnterpriseDB Published 10 days ago

EDB Postgres for Kubernetes is a certified Level 5 Operator for Red Hat OpenShift

- ▶ This is designed to streamline Day 2 operations of PostgreSQL databases
- ▶ Enhanced Database Management
- ▶ Supports point-in-time recovery (PITR)
- ▶ Ensures robust data protection and recovery options
- ▶ Integration with business continuity solutions such as Red Hat OpenShift API for Data Protection (OADP) and Veeam Kasten, Trilio, Portworx Backup, IBM Fusion, and others

Proven Solution at Scale



EDB Postgres on OpenShift use cases

- Cloud-Native Database Deployment
- Database as a Service (DBaaS)
- High Availability and Disaster Recovery (HA & DR)
- DevOps and Continuous Integration/Continuous Deployment (CI/CD)
- Microservices and Application Modernization
- Move from VMWare to OpenShift
- Data Security and Compliance (using TDE and Advanced Security provided by EPAS)
- Hybrid and Multi-Cloud Deployments
- Multi-Tenant Applications (isolation)
- AI/ML Sovereign Platform

Euro Information

Company profile

Euro-Information is the fintech company of the Crédit Mutuel group. Euro-Information manages the IT systems of 16 federations of Crédit Mutuel as well as those of CIC and of all the financial, insurance, property, consumer credit, private banking, financing, telephony and technological subsidiaries.



- Red Hat OpenShift
- EDB Postgres for Kubernetes
- PostgreSQL
- EPAS



- EDB considerably reduces IT costs associated with database maintenance.
- 280 cores: Enterprise Plan + Production Support

Summary

Use Case

On prem DBaaS (**in Production**)

Workload

Transactional

Application Name

All internal Postgres applications

EDB Tools of Interest

PostgreSQL and EDB Postgres for Kubernetes

Problem

- Fast database deployment
- Adopt a supported and secure Open Source platform
- Onprem DBaaS
- Align to in-house RDBMS standardization

Solution

- Use Postgres capabilities to build and maintain local applications
- Use Red Hat OpenShift platform to accelerate the provisioning of databases and applications

Results

- Applications running with PostgreSQL databases in a centralized environment
- Massive reduction of TCO of database service operations

Airbus

Company profile

Airbus SE is a European aerospace corporation. The company's primary business is the design and manufacturing of commercial aircraft but it also has separate defence and space and helicopter divisions.



- Red Hat OpenShift
- EDB Postgres for Kubernetes
- EPAS
- TDE



- Improve database deployment speed
- Reduce DB support
- Cost reduction

Summary

Use Case

On prem DBaaS (Production)

Workload

Transactional

Application Name

All VMWare PostgreSQL databases

EDB Tools of Interest

EDB Postgres Advanced Server with Oracle and TDE (optional)

Problem

- Flexibility
- Cost reduction
- New managed service in OpenShift

Solution

- EDB Postgres for Kubernetes with EPAS. Depending of the applications needs, EPAS and/or TDE will be activated

Results

- POC done
- Decision is taken
- Number of cores not yet communicated

Banque de France

Company profile

The Banque de France is France's central bank. A two-hundred-year-old institution, privately-owned when it was founded on January 18, 1800 under the Consulate by General Bonaparte, it became state-owned on January 1, 1946 when it was nationalized by General de Gaulle.



- Red Hat OpenShift
- CloudNativePG
- PostgreSQL



- 100 cores
- Subscription plan:
 - Community360 plan + Production Support

Summary

Use Case

OnPrem DBaaS (in production)

Workload

Transactional

Application Name

Multiple applications

EDB Tools of Interest

PostgreSQL, CloudNativePG

Problem

- Fast database deployment
- Provide containerized Postgres DBaaS

Solution

- Use OpenShift to provide this service with the operator
- Fast deployment and with Open Source database

Results

- OpenShift based PostgreSQL cluster deployments expand the internal offering alongside traditional VM based database cluster deployments



La Poste

Company profile

La Poste is a postal service company in France, operating in Metropolitan France, the five French overseas departments and regions and the overseas collectivity of Saint Pierre and Miquelon. Under bilateral agreements, La Poste also has responsibility for mail services in Monaco through La Poste Monaco and in Andorra alongside the Spanish company Correos.



- Red Hat OpenShift
- EDB Postgres for Kubernetes
- PostgreSQL



- EDB considerably reduces IT costs associated with database maintenance.
- 12 Cores: Standard Plan + Premium Support

Summary

Use Case

On prem DBaaS with HA and DR
(In Production)

Workload

Transactional

Application Name

Portail XaaS

EDB Tools of Interest

PostgreSQL and EDB Postgres for Kubernetes

Problem

- Provide a database HA solution for Ansible Automation Platform (AAP)
- Database must be in HA and DR

Solution

- Use EDB Postgres for Kubernetes to provide a HA and DR solution for PostgreSQL databases
- Deploy in 2 OpenShift clusters our operator

Results

- La Poste developer can use their internal 'La Post Service Portal' to provision more than 64 backends.
- Reduce risk deploying EDB solutions.

Thank you

Red Hat is the world's leading provider of enterprise open source software solutions. Award-winning support, training, and consulting services make Red Hat a trusted adviser to the Fortune 500.



linkedin.com/company/red-hat



youtube.com/user/RedHatVideos



facebook.com/redhatinc




twitter.com/RedHat

Introduction to Postgres and EDB



PostgreSQL is Booming !




35% 

Thirty-five percent of Enterprises will consider Postgres for their next project



EnterpriseDB Biggest PostgreSQL Contributor



CLIENTS

AGENTS

APPLICATIONS

USERS

**EDB POSTGRES[®]
AI**

WORKLOADS

OPERATIONAL

RELATIONAL

DOCUMENT

VECTOR

GRAPH

OPERATIONAL

ANALYTICAL

WAREHOUSE

LAKEHOUSE

REAL-TIME

BATCH

AGENTIC

RUNTIME

BUILDING

SEMANTIC SEARCH

RAG

MACHINE LEARNING

PLATFORM SERVICES

GOVERNANCE

SECURITY

DATA INTEGRATION

MIGRATIONS

OBSERVABILITY

LIFECYCLE

HA/DR

DATA LAYER

LAKEHOUSE
OBJECT STORAGE

DATABASES
POSTGRESQL

STREAMS
EVENT DATA

DEPLOYMENT OPTIONS

IBM MAINFRAME

SOVEREIGN DC

PUBLIC CLOUD

NVIDIA GPU

EDB AI FACTORY HW



CNPG Operator: Reference Architecture and functionalities



Kubernetes timeline

- 2014, June: Google open sources Kubernetes
- 2015, July: Version 1.0 is released
- 2015, July: Google and Linux Foundation start the CNCF
- 2016, November: The operator pattern is introduced in a blog post
- 2018, August: The Community takes the lead
- 2019, April: Version 1.14 introduces **Local Persistent Volumes**
- 2019, August: EDB team starts the Kubernetes initiative
- 2020, June: we publish this blog about benchmarking local PVs on bare metal
- 2020, June: Data on Kubernetes Community founded
- 2021, February: EDB Cloud Native Postgres (CNP) 1.0 released
- 2022, May: **EDB donates CNP** and open sources it under CloudNativePG
- 2025, January: CloudNativePG was recognized as an official **#CNCF** project



A kubernetes operator for Postgres



Kubernetes adoption is rising and it is already the de facto **standard orchestration tool**



PostgreSQL clusters “**management the kubernetes way**” enables many cloud native usage patterns, e.g. spinning up, disposable clusters during tests, one cluster per microservice and one database per cluster



CNPG tries to encode years of experience managing PostgreSQL clusters into **an Operator which should automate all the known tasks a user could be willing to do**

Our PostgreSQL operator must simulate a part of the work of a DBA



Autopilot

It automates the steps that a human operator would do to deploy and to manage a Postgres database inside Kubernetes, including automated failover.



Security

A man in a tactical vest with the word "SECURITY" printed on the back, holding a walkie-talkie to his mouth. The image is faded and serves as a background for the slide.

SECURITY

CloudNativePG is secured by default.



It doesn't rely on statefulsets and uses its own way to manage persistent volume claims where the PGDATA is stored.

Data persistence



Designed for Kubernetes

It's entirely declarative, and directly integrates with the Kubernetes API server to update the state of the cluster — for this reason, it does not require an external failover management tool.



Features

Deployment	Administration	Backup & Recovery	Monitoring	Security	High Availability
Kubernetes operator	Single node	Backup	Prometheus	TDE	Switchover
Kubernetes plugin	Cluster (Multi node)	Recovery	Grafana dashboards	Certificates	Failover
EDB Postgres (EPAS)	PostgreSQL configuration	PITR	Postgres Enterprise Manager	Data redaction	Scale out / scale down
PostGIS	Pooling	Volume Snapshots	Logging	Password management	Minor / Major updates




Use Cases



Use case 1 architecture

A single database is the simplest setup, involving one instance of a database server.


- Development and testing environments
- Small applications with low traffic
- Non-critical data analysis
- Applications with high tolerance for downtime
- Cost-sensitive projects



Use case 2 architecture

An HA database setup aims to minimize downtime by having redundant components. If one component fails, another takes over automatically or with minimal intervention. This usually involves techniques like clustering, replication, or mirroring **within the same data center or availability zone**.


- Business critical Applications
- Applications with stringent SLAs
- Real-time systems
- Improving user experience
- Minimizing planned downtime




Use case 3 architecture

A DR database setup focuses on protecting data and ensuring business continuity in the event of a large-scale disaster affecting an entire data center or region (e.g., natural disasters, power outages, cyberattacks). This typically involves replicating data to a geographically separate location.

- Regulatory compliance
- Protecting against catastrophic data loss
- Ensuring business continuity for mission-critical systems



Use case 3 architecture



Interactive session
It's time to go hands-on!

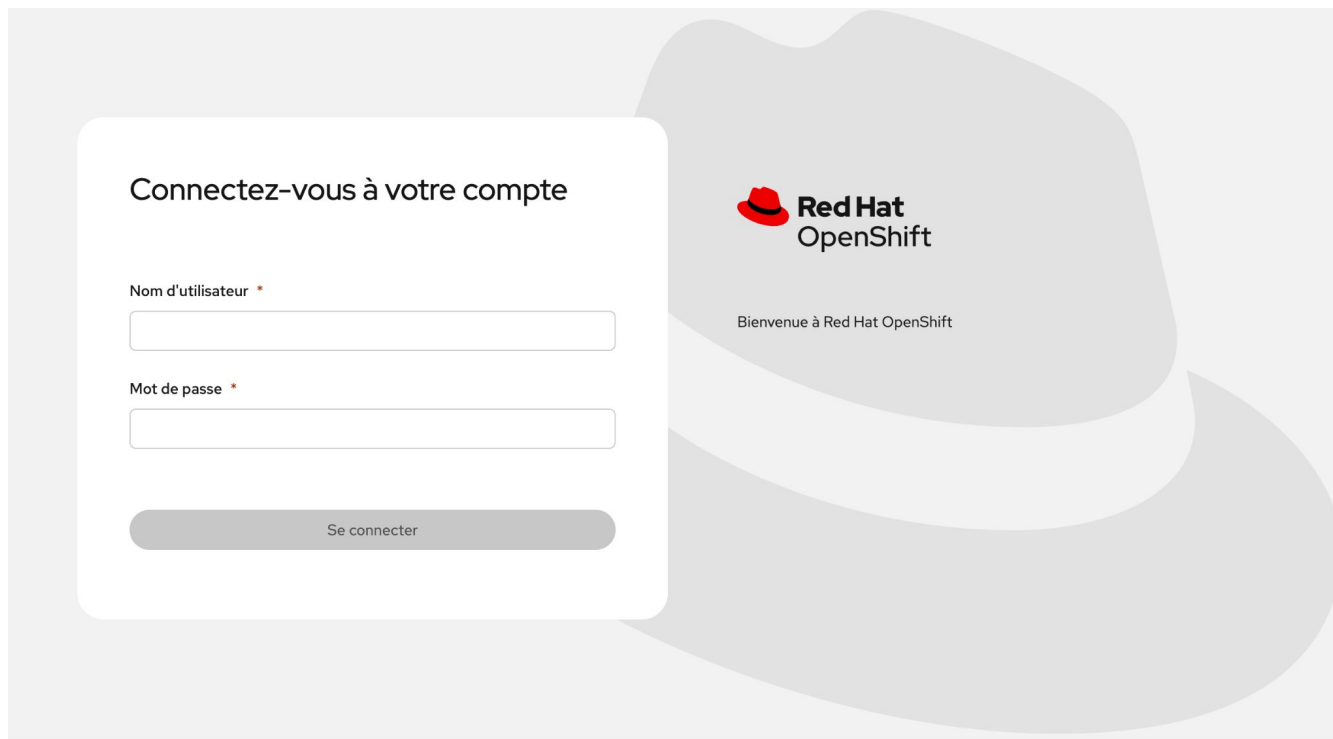


Environment

OpenShift Console	https://console-openshift-console.apps.cluster-dkwbs.dkwbs.sandbox2953.opentlc.com/
OpenShift User name	user01..user30
OpenShift Password	edbredhat
Minio UI	https://minio-ui-minio.apps.cluster-dkwbs.dkwbs.sandbox2953.opentlc.com/
Minio internal service	http://minio-service.minio.svc.cluster.local:9000
Minio API	https://minio-api-minio.apps.cluster-dkwbs.dkwbs.sandbox2953.opentlc.com
Minio User	minio
Minio password	edb-workshop



Red Hat OpenShift Login



The image shows a login form for Red Hat OpenShift. The form is white with rounded corners and is set against a light gray background featuring a large, faint silhouette of a Red Hat. The form contains the following elements:

- Header:** "Connectez-vous à votre compte" (Log in to your account).
- Username Field:** Labeled "Nom d'utilisateur *" (Username *), with a red asterisk indicating a required field. Below the label is a white input box.
- Password Field:** Labeled "Mot de passe *" (Password *), with a red asterisk indicating a required field. Below the label is a white input box.
- Login Button:** A gray button with rounded ends, labeled "Se connecter" (Log in).
- Red Hat OpenShift Logo:** Located to the right of the form, consisting of a red hat icon and the text "Red Hat OpenShift".
- Welcome Text:** Below the logo, the text "Bienvenue à Red Hat OpenShift" (Welcome to Red Hat OpenShift) is displayed.

Red Hat OpenShift environment

The screenshot shows the Red Hat OpenShift console interface. On the left is a sidebar with navigation links: Home, Favorites, Ecosystem, Helm, Workloads, Networking, Storage, Builds, Observe, Compute, User Management, and Administration. The main content area displays the 'Installed Operators' for the 'openshift-image-registry' project. A blue banner at the top of the main area states: 'You are logged in as a temporary administrative user. Update the [cluster OAuth configuration](#) to allow others to log in.' Below this, the 'Installed Operators' section includes a search bar and a table of installed operators. The table has columns for Name, Managed Namespaces, Status, and Provided APIs. One operator is listed: 'EDB Postgres for Kubernetes', which is provided by EDB and is running in the 'openshift-image-registry' namespace. The status is 'Succeeded' and 'Up to date'. The provided APIs include Backups, Cluster Image Catalog, Cluster, Postgres Database, and a link to view more.

Red Hat OpenShift

You are logged in as a temporary administrative user. Update the [cluster OAuth configuration](#) to allow others to log in.

Project: openshift-image-registry

Installed Operators

Installed Operators are represented by ClusterServiceVersions within this Namespace. For more information, see the [Understanding Operators documentation](#). Or create an Operator and ClusterServiceVersion using the [Operator SDK](#).

Name Search by name...

Name	Managed Namespaces	Status	Provided APIs
EDB Postgres for Kubernetes 1.28.0 provided by EDB	openshift-image-registry The operator is running in openshift-operators but is managing this namespace	✓ Succeeded Up to date	Backups Cluster Image Catalog Cluster Postgres Database View 6 more...



Red Hat OpenShift operator

The screenshot displays the Red Hat OpenShift console interface. On the left is a sidebar with navigation links: Home, Favorites, Ecosystem, Software Catalog, Installed Operators, Helm, Workloads, Networking, Storage, Builds, Observe, Compute, User Management, and Administration. The main content area shows the details for the 'EDB Postgres for Kubernetes' operator, version 128.0, provided by EDB. A top banner indicates the user is logged in as a temporary administrative user and suggests updating the cluster OAuth configuration. Below this, a breadcrumb trail shows 'Installed Operators' > 'Operator details'. The operator name 'EDB Postgres for Kubernetes' is displayed with a star icon and an 'Actions' dropdown. A horizontal tab bar includes links for Details (selected), YAML, Subscription, Events, All instances, Backups, Cluster Image Catalog, Cluster, Postgres Database, Failover Quorum, and Image Catalog. The 'Provided APIs' section contains nine cards, each with an icon, title, description, and a 'Create instance' link. The right sidebar provides metadata: Provider (EDB), Created at (Dec 17, 2025, 10:38 AM), Links (including EDB Postgres for Kubernetes product and documentation URLs), and Maintainers (Jonathan Gonzalez V, Jonathan Battiato, and Niccolo Fei).

You are logged in as a temporary administrative user. Update the [cluster OAuth configuration](#) to allow others to log in.

Project: openshift-image-registry

[Installed Operators](#) > Operator details

EDB Postgres for Kubernetes 128.0 provided by EDB Actions

< Details YAML Subscription Events All instances Backups Cluster Image Catalog Cluster Postgres Database Failover Quorum Image Catalog >

Provided APIs

Icon	API Name	Description	Create Instance Link
	Backups	PostgreSQL backup (physical base backup)	Create instance
	Cluster Image Catalog	A cluster-wide catalog of PostgreSQL operand images	Create instance
	Cluster	PostgreSQL cluster (primary/standby architecture)	Create instance
	Postgres Database	Declarative creation and management of a database on a Cluster	Create instance
	Failover Quorum	FailoverQuorum contains the information about the current failover quorum status of a PG cluster	Create instance
	Image Catalog	A catalog of PostgreSQL operand images	Create instance
	Pooler	Pooler for a Postgres Cluster (with PgBouncer)	Create instance
	Postgres Publication	Declarative creation and management of a Logical Replication Publication in a PostgreSQL Cluster	Create instance
	Scheduled Backups	Backup scheduler for a given Postgres cluster	Create instance

Provider
EDB

Created at
Dec 17, 2025, 10:38 AM

Links
EDB Postgres for Kubernetes
<https://www.enterprisedb.com/products/postgres-ql-on-kubernetes-ha-clusters-k8s-containers-scalable>
Documentation
<https://www.enterprisedb.com/docs/postgres-for-kubernetes/latest/>

Maintainers
Jonathan Gonzalez V. jonathan.gonzalez@enterprisedb.com
Jonathan Battiato jonathan.battiato@enterprisedb.com
Niccolo Fei niccolo.fei@enterprisedb.com
Gabriele Bartolini gabriele.bartolini@enterprisedb.com



Use case

The environment



Features shown during the demo

- Kubernetes plugin install
- Check the CloudNativePG operator status
- Postgres cluster install
- Insert data in the cluster
- Failover
- Backup
- Recovery
- Scale out/down
- Fencing
- Hibernation
- Monitoring
- Rolling updates (minor and major)

Deployment

Administration

Backup and
Recovery

High Availability

Monitoring

Last CloudNativePG tested version is 1.25



This demo is in 

[https://github.com/raphael-chir/edb-postgres-f
or-kubernetes-in-openshift](https://github.com/raphael-chir/edb-postgres-f
or-kubernetes-in-openshift)






Hands-on



OpenShift and DevSpaces access

<https://github.com/raphael-chir/edb-postgres-for-kubernetes-in-openshift>



Plugging and OpenShift login

Scripts


01_install_plugin.sh

Install plugin

CloudNativePG provides a plugin for kubectl to manage a Postgres cluster in Kubernetes. You can manage status, promote, certificates, restart, reload, maintenance, report, logs, destroy, hibernation, backups, ...

Doc [link](#)

OpenShift Login and namespace creatop,



Code

```
# Install plugin

curl -sSfL \
  https://github.com/EnterpriseDB/kubectl-cnp/raw/main/install.sh
| sh -s -- -b .

# Login to the OpenShift Cluster

oc login ...

# Create your namespace

oc new-project userXX
```

Checking operator installation

Scripts

03_check_operator_installed.sh

Checking operator installation

In Kubernetes, the operator is by default installed in the postgresql-operator-system namespace as a Kubernetes Deployment. The name of this deployment depends on the installation method. When installed through the manifest or the cnp plugin, by default, it is called postgresql-operator-controller-manager. When installed via Helm, by default, the deployment name is derived from the helm release name, appended with the suffix -edb-postgres-for-kubernetes (e.g., <name>-edb-postgres-for-kubernetes).

Doc [link](#)

Code

```
# Check operator installation

oc get deploy -A | grep postgres
NAMESPACE:    openshift-operators
NAME:         postgresql-operator-controller-manager
READY:        1/1
UP-TO-DATE:   1
AVAILABLE:    1
AGE:          5d21h
```



Install Postgres cluster

Scripts


```
04-prepare-cluster.sh
05_install_cluster.sh
06_show_status.sh
07_insert_data.sh
```

Prepare Postgres cluster template

Setup credentials (AWS or Minio) and prepare context for your users

Deploy Postgres cluster in HA

Different types of architectures can be deployed. In our demo, 3 instances are deployed with HA available.



Insert Data

Insert data in test table

Doc [link](#)

Code

```
export cluster_name="cluster-userXX"
export TMP=$TMPDIR

# Get template
envsubst < templates/${cluster_name}-template.yaml >
$TMP/${cluster_name}.yaml

# Deploy Postgres cluster
${kubectl_cmd} apply -f $TMP/${cluster_name}.yaml

-----
ApiVersion: postgresql.k8s.enterprisedb.io/v1
kind: Cluster
metadata:
  name: cluster-example
spec:
  instances: 3
  imageName: "docker.enterprisedb.com/k8s/edb-postgres-advanced:17.6"
  enableSuperuserAccess: true

replicationSlots:
  highAvailability:
    enabled: true
```



Monitoring

Show status

Cluster Summary

Name: default/cluster-sergio-1
System ID: 7566239054312464401
PostgreSQL Image: docker.enterprisedb.com/k8s/edb-postgres-advanced:17.6
Primary instance: cluster-sergio-1
Primary promotion time: 2025-10-28 11:52:15 +0000 UTC (4m30s)
Status: Cluster in healthy state
Instances: 3
Ready instances: 3
Size: 127M
Current Write LSN: 0/6032300 (Timeline: 1 - WAL File: 00000001000000000000000006)

Continuous Backup status (Barman Cloud Plugin)

ObjectStore / Server name: object-store/cluster-sergio
First Point of Recoverability: -
Last Successful Backup: -
Last Failed Backup: -
Working WAL archiving: OK
WALs waiting to be archived: 0
Last Archived WAL: 000000010000000000000005.00000028.backup @ 2025-10-28T11:54:11.908195Z
Last Failed WAL: -

Streaming Replication status

Replication Slots Enabled											
Name	Sent LSN	Write LSN	Flush LSN	Replay LSN	Write Lag	Flush Lag	Replay Lag	State	Sync State	Sync Priority	Replication Slot
cluster-sergio-2	0/9000000	0/9000000	0/9000000	0/9000000	00:00:00	00:00:00	00:00:00	streaming	quorum	1	active
cluster-sergio-3	0/9000000	0/9000000	0/9000000	0/9000000	00:00:00	00:00:00	00:00:00	streaming	quorum	1	active

Instances status

Name	Current LSN	Replication role	Status	QoS	Manager Version	Node
cluster-sergio-1	0/9000000	Primary	OK	Guaranteed	1.28.0	ocp4-multiarch-worker1-ppc64
cluster-sergio-2	0/9000000	Standby (sync)	OK	Guaranteed	1.28.0	ocp4-multiarch-w5gg4-master-2
cluster-sergio-3	0/9000000	Standby (sync)	OK	Guaranteed	1.28.0	ocp4-multiarch-w5gg4-master-1



Promote and Failover

Scripts

08_promote.sh

Promote

The meaning of this command is to promote a pod in the cluster to primary, so you can start with maintenance work or test a switch-over situation in your cluster:

Doc [link](#)

Code

```
# Promote
. ./config.sh

${kubectl_cnp} promote ${cluster_name} ${cluster-name}-2
```



Minor upgrade

Scripts

09_upgrade.sh

Minor upgrade

To upgrade to a newer minor version, simply update the PostgreSQL container image reference in your cluster definition, either directly or via image catalogs. CloudNativePG will trigger a rolling update of the cluster, replacing each instance one by one, starting with the replicas. Once all replicas have been updated, it will perform either a switchover or a restart of the primary to complete the process.

Doc [link](#)

Code

```
# Minor upgrade

# From
apiVersion: postgresql.k8s.enterprisedb.io/v1
kind: Cluster
metadata:
  name: cluster-sergio1
spec:
  instances: 3
  imageName:
    "docker.enterprisedb.com/k8s/edb-postgres-advanced:17.6"
  ...

# To
apiVersion: postgresql.k8s.enterprisedb.io/v1
kind: Cluster
metadata:
  name: cluster-sergio1
spec:
  instances: 3
  imageName:
    "docker.enterprisedb.com/k8s/edb-postgres-advanced:17.7"
  ...
```



Backups

Scripts

```
10_backup_cluster.sh  
11_backup_describe.sh
```

Backup

Backups are used using Barman cloud with the new Barman Plugin. In case the plugin is not compatible (like in Red Hat OpenShift), we will continue using Barman Cloud (without the plugin).

PostgreSQL natively provides first class backup and recovery capabilities based on file system level (physical) copy. These have been successfully used for more than 15 years in mission critical production databases, helping organizations all over the world achieve their disaster recovery goals with Postgres.

Doc [link](#)

Code

```
cat $TMP/backup.yaml  
  
apiVersion: postgresql.k8s.enterprisedb.io/v1  
kind: Backup  
metadata:  
  name: cluster-example-backup-test  
spec:  
  cluster:  
    name: cluster-example  
  
${kubect1_cmd} apply -f ./yaml/backup.yaml
```



Recovery

Scripts

12_restore_cluster.sh

Recovery

In PostgreSQL, recovery refers to the process of starting an instance from an existing physical backup. PostgreSQL's recovery system is robust and feature-rich, supporting Point-In-Time Recovery (PITR)—the ability to restore a cluster to any specific moment, from the earliest available backup to the latest archived WAL file.

Doc [link](#)

Code

```
cat $TMP/restore.yaml

apiVersion: postgresql.k8s.enterprisedb.io/v1
kind: Cluster
metadata:
  name: cluster-restore
spec:
  instances: 1
  imageName: quay.io/enterprisedb/postgresql:16.4
  imagePullPolicy: IfNotPresent
  bootstrap:
    recovery:
      source: source
  plugins:
  - name: barman-cloud.cloudnative-pg.io
    isWALArchiver: true
    parameters:
      barmanObjectName: object-store
  externalClusters:
  - name: source
    plugin:
      name: barman-cloud.cloudnative-pg.io
      parameters:
        barmanObjectName: object-store
        serverName: cluster-example

${kubectl_cmd} apply -f $TMP/restore.yaml
```



Promote and Failover

Scripts

13_failover.sh

Automated Failover

In the case of unexpected errors on the primary for longer than the `.spec.failoverDelay` (by default 0 seconds), the cluster will go into failover mode. This may happen, for example, when:

- The primary pod has a disk failure
- The primary pod is deleted
- The postgres container on the primary has any kind of sustained failure

In the failover scenario, the primary cannot be assumed to be working properly.
Doc [link](#)

Code

```
# Automatic Failover (simulate failover killing pods of primary)
${kubect1_cmd} delete pvc/${primary} \
                      pvc/${primary}-wal \
                      pod/${primary} \
                      --force
```



Scale out, scale down

Scripts

```
14_scale_out.sh  
15_scale_down.sh
```

Scale out and down

The operator allows you to scale up and down the number of instances in a PostgreSQL cluster. New replicas are started up from the primary server and participate in the cluster's HA infrastructure.

Doc [link](#)

Code

```
# Scale out  
kubectl scale cluster cluster-example --replicas=4  
  
# Scale down  
kubectl scale cluster cluster-example --replicas=2
```



Fencing and Hibernation

Scripts

```
30_fencing_on.sh  
31_fencing_off.sh
```

Fencing

Fencing is the process of protecting the data in one, more, or even all instances of a PostgreSQL cluster when they appear to be malfunctioning. When an instance is fenced, the PostgreSQL server process is guaranteed to be shut down, while the pod is kept running. This ensures that, until the fence is lifted, data on the pod isn't modified by PostgreSQL and that you can investigate file system for debugging and troubleshooting purposes.

Doc [link](#)

Code

```
# Fencing on  
${kubectl_cnp} fencing on ${cluster_name} ${replica}  
  
# Fencing off  
${kubectl_cnp} fencing off ${cluster_name} ${replica}
```



Fencing and Hibernation

Scripts

32_hibernation_on.sh
33_hibernation_off.sh

Hibernation

CloudNativePG supports hibernation of a running PostgreSQL cluster in a declarative manner, through the `cnpg.io/hibernation` annotation. Hibernation enables saving CPU power by removing the database pods while keeping the database PVCs. This feature simulates scaling to 0 instances.

Doc [link](#)

Code

```
# Hibernation on
${kubectl_cmd} annotate cluster ${cluster_name} \
    --overwrite k8s.enterprisedb.io/hibernation=on

# Hibernation off
${kubectl_cmd} annotate cluster ${cluster_name} \
    --overwrite k8s.enterprisedb.io/hibernation=off
```



Major upgrade (copying data)

Scripts

20_upgrade_major_version.sh

Major upgrade (with native logical replication)

Major PostgreSQL releases introduce changes to the internal data storage format, requiring a more structured upgrade process.

CloudNativePG supports three methods for performing major upgrades:

- Logical dump/restore – Blue/green deployment, offline.
- Native logical replication – Blue/green deployment, online.
- **Physical with pg_upgrade** – In-place upgrade, offline (covered in the "Offline In-Place Major Upgrades" section below).

Doc [link](#)

Code

```
apiVersion: postgresql.k8s.enterprisedb.io/v1
kind: Cluster
metadata:
  name: ${cluster_major_upgrade}
spec:
  instances: 1
  imageName: ${postgres_major_upgrade_image}

  storage:
    size: 2Gi

  bootstrap:
    initdb:
      import:
        type: microservice
        databases:
          - app
        source:
          externalCluster: ${cluster_name}

  externalClusters:
    - name: ${cluster_name}
      connectionParameters:
        # Use the correct IP or host name for the source database
        host: ${cluster_name}-rw
        user: postgres
        dbname: postgres
      password:
        name: ${cluster_name}-superuser
        key: password
```



Major upgrade (copying data)

Scripts


20_upgrade_major_version.sh

Major upgrade (with native logical replication)

- In this step we will migrate the app database from our existing cluster to the new cluster
- We will create the yaml file with the setting **"import"** in the bootstrap section
- The operator uses internally postgres tools `pg_dump` and `pg_restore`
- This method can be used to **migrate** another database or to run

out-of-the-place upgrade

- Possible settings:
- Microservices
- Monolith



Code

```
apiVersion: postgresql.k8s.enterprisedb.io/v1
kind: Cluster
metadata:
  name: cluster-sergio-major
spec:
  instances: 1
  imageName: docker.enterprisedb.com/k8s/postgresql:18.1
  imagePullSecrets:
    - name: postgresql-operator-pull-secret
  enableSuperuserAccess: true

  bootstrap:
    initdb:
      import:
        type: microservice
        databases:
          - postgres
        source:
          externalCluster: cluster-sergio

  externalClusters:
    - name: cluster-sergio
      connectionParameters:
        host: cluster-sergio-rw
        user: postgres
        dbname: postgres
      password:
        name: cluster-sergio-superuser
        key: password
    ...
```

Already finished ?



Application deployment

Database : CNPG Potsgres cluster
Backend : Rest API - Postgrest
Frontend : Swagger UI
Frontend : World Cup UI

Update manifest postgres-secret.yaml

Get the URI of your instance

```
oc get secret cluster-<NAMESPACE>-app -o jsonpath='{.data.uri}' |  
base64 -d  
Replace in PGRST_DB_URI
```

Update manifest postgres-deployment.yaml

Adapt the fqdn

```
PGRST_OPENAPI_SERVER_PROXY_URI =  
http://postgrest-<namespace>.apps-crc.testing
```

```
PGRST_SERVER_CORS_ALLOWED_ORIGINS  
http://swagger-ui-<namespace>.apps-crc.testing,http://worldcup-ui-<name  
space>.apps-crc.testing
```

Update the frontend to point to your backend

Deploy `./packaging/deploy.sh`

Code

▼ edb-postgres-for-kubernetes-in-openshift

- ▼ app
 - > database
 - ▼ packaging
 - > db
 - ▼ k8s
 - ! db-dump-configmap.yaml U
 - ! db-restore-job.yaml U
 - ! frontend.yaml U
 - ! postgrest-config.yaml U
 - ! postgrest-deployment.yaml U
 - ! postgrest-secret.yaml U
 - ! postgrest-service-route.yaml U
 - ! swagger-ui.yaml U
 - \$ deploy.sh U
 - i README.md U
 - \$ undeploy.sh U
 - > worldcup-ui
 - i README.md U



Thank you

For any queries, contact:

sergio.romera@enterprisedb.com



Backup Slides



Architectures



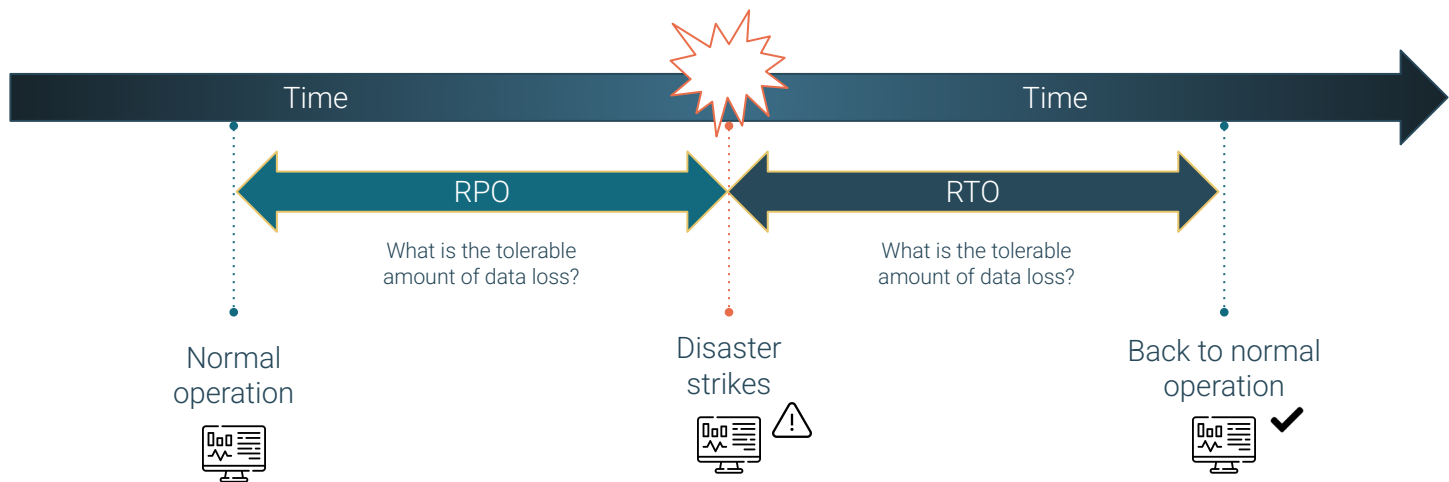
When to choose Kubernetes over VMs?

- 01 |** Cloud Native Applications that already run in Kubernetes
- 02 |** Scalable, replicated databases
- 03 |** Applications requiring automated failover and self-healing
- 04 |** Teams skilled in Kubernetes who want a unified infrastructure




Concepts

- Recovery Point Objective (**RPO**) and Recovery Time Objective (**RTO**) are key concepts in disaster recovery and business continuity planning, particularly related to data loss and system downtime.




Red Hat Recommendation

Red Hat recommend stretched clusters ONLY when latencies don't exceed 5 milliseconds (ms) round-trip time (RTT) between the nodes in different locations, with a maximum RTT of 10 ms.




Two separate single data center Kubernetes clusters


In case you cannot go beyond two data centers and you end up with two separate Kubernetes clusters, don't despair.




Shared workload, shared storage




Shared workload, shared storage



Shared workload, shared storage




Shared workloads, local storage



Recommended architectures

<https://www.cncf.io/blog/2023/09/29/recommended-architectures-for-postgresql-in-kubernetes/>




PostgreSQL most loved database (2023)

Source: [Stackoverflow](#)

This year, PostgreSQL took over the first place spot from MySQL. Professional Developers are more likely than those learning to code to use PostgreSQL (50%) and those learning are more likely to use MySQL (54%).

MongoDB is used by a similar percentage of both Professional Developers and those learning to code and it's the second most popular database for those learning to code (behind MySQL).



Links:

Openshift Console:

<https://console-openshift-console.apps.cluster-826x2.826x2.sandbox5464.opentlc.com/>

Users:

name: user2..user32
Password: edb-workshop

Devspaces url:

<https://devspaces.apps.cluster-826x2.826x2.sandbox5464.opentlc.com>

Short url to Devspaces:

<https://tinyurl.com/mrxs95pb>

Minio:

UI: <https://minio-ui-minio.apps.cluster-826x2.826x2.sandbox5464.opentlc.com>
API: <https://minio-api-minio.apps.cluster-826x2.826x2.sandbox5464.opentlc.com>


User: minio
Password: edb-workshop




Database Migration

- In this step we will migrate the app database from our existing cluster to the new cluster
- We will create the yaml file with the setting “**import**” in the bootstrap section
- The operator uses internally postgres tools pg_dump and pg_restore
- This method can be used to **migrate** another database or to run **out-of-the-place upgrade**
- Possible settings:

Microservices:



Monolith:



Major upgrade (copying data)

Scripts

```
24_major_upgrade_in_place.sh  
25_verify_major_upgrade_16_17.sh
```

Major upgrade (offline In-Place Major Upgrades)

CloudNativePG performs an offline in-place major upgrade when a new operand container image with a higher PostgreSQL major version is declaratively requested for a cluster.

Major upgrades are only supported between images based on the same operating system distribution. For example, if your previous version uses a bullseye image, you cannot upgrade to a bookworm image.

Doc [link](#)

Code

```
apiVersion: postgresql.k8s.enterprisedb.io/v1  
kind: Cluster  
metadata:  
  name: ${cluster_major_upgrade}  
spec:  
  instances: 1  
  imageName: ${postgres_major_upgrade_image}  
  
  storage:  
    size: 2Gi  
  
  bootstrap:  
    initdb:  
      import:  
        type: microservice  
        databases:  
          - app  
  ...
```



Disaster recovery

Scripts

```
./dr/01-create-all.sh
./dr/02-insert-eu.sh
./dr/03-switchover-cluster.sh
```

Disaster recovery


A replica cluster is a CloudNativePG Cluster resource designed to replicate data from another PostgreSQL instance, ideally also managed by CloudNativePG.

Typically, a replica cluster is deployed in a different Kubernetes cluster in another region. These clusters can be configured to perform cascading replication and can rely on object stores for data replication from the source, as detailed further down..

Two clusters:

- pg-eu
- pg-us

Doc [link](#)



Code

```
# EU
kubectl delete objectstores.barmancloud.cnpg.io object-storage-eu
envsubst < ../templates/dr/azure-object-storage-eu.yaml >
./tmp/object-storage-eu.yaml
kubectl apply -f ./tmp/object-storage-eu.yaml
sleep 2
rm -f ./tmp/pg-eu.yaml
envsubst < ../templates/dr/pg-eu.yaml > ./tmp/pg-eu.yaml
print info "Creating pg-eu cluster...\n"
kubectl apply -f ./tmp/pg-eu.yaml
print_info "pg-eu cluster created.\n"

# US
kubectl delete objectstores.barmancloud.cnpg.io object-storage-us
envsubst < ../templates/dr/azure-object-storage-us.yaml >
./tmp/object-storage-us.yaml
kubectl apply -f ./tmp/object-storage-us.yaml
sleep 2
rm -f ./tmp/pg-us.yaml
envsubst < ../templates/dr/pg-us.yaml > ./tmp/pg-us.yaml
kubectl wait --timeout=30m --for=condition=Ready cluster/pg-eu
print info "Creating pg-us cluster...\n"
kubectl apply -f ./tmp/pg-us.yaml
./backup_cluster.sh pg-eu
print_info "pg-us cluster created.\n"
```

The “cnp” plugin for kubectl

- The official CLI for CloudNativePG
 - Available also as RPM or Deb package
- Extends the ‘kubectl’ command:
 - Customize the installation of the operator
 - Status of a cluster
 - Perform a manual switchover (promote a standby) or a restart of a node
 - Issue TLS certificates for client authentication
 - Declare start and stop of a Kubernetes node maintenance
 - Destroy a cluster and all its PVC
 - Fence a cluster or a set of the instances
 - Hibernate a cluster
 - Generate jobs for benchmarking via pgbench and fio
 - Issue a new backup
 - Start pgadmin



Name: cluster-example
Namespace: default
System ID: 7100921006673293335
PostgreSQL Image: ghcr.io/cloudnative-pg/postgresql:14.3
Primary instance: cluster-example-2
Status: Cluster in healthy state
Instances: 3
Ready instances: 3
Current Write LSN: 0/C000060 (Timeline: 4 - WAL File: 00000004000000000000000C)

Certificates Status

Certificate Name	Expiration Date	Days Left Until Expiration
cluster-example-replication	2022-08-21 13:15:00 +0000 UTC	89.95
cluster-example-server	2022-08-21 13:15:00 +0000 UTC	89.95
cluster-example-ca	2022-08-21 13:15:00 +0000 UTC	89.95

Continuous Backup status

First Point of Recoverability: 2022-05-23T13:37:08Z
Working WAL archiving: OK
WALs waiting to be archived: 0
Last Archived WAL: 00000004000000000000000B @ 2022-05-23T13:42:09.37537Z
Last Failed WAL: -

Streaming Replication status

Name	Sent LSN	Write LSN	Flush LSN	Replay LSN	Write Lag	Flush Lag	Replay Lag	State	Sync State	Sync Priority
cluster-example-3	0/C000060	0/C000060	0/C000060	0/C000060	00:00:00	00:00:00	00:00:00	streaming	async	0
cluster-example-1	0/C000060	0/C000060	0/C000060	0/C000060	00:00:00	00:00:00	00:00:00	streaming	async	0

Instances status

Name	Database Size	Current LSN	Replication role	Status	QoS	Manager Version
cluster-example-3	33 MB	0/C000060	Standby (async)	OK	BestEffort	1.15.0
cluster-example-2	33 MB	0/C000060	Primary	OK	BestEffort	1.15.0
cluster-example-1	33 MB	0/C000060	Standby (async)	OK	BestEffort	1.15.0



Install CNPG plugin

NOT NEEDED DURING WORKSHOP
For illustrative purposes.

- In the web terminal run the script 01_install_plugin.sh:
./01_install_plugin.sh



Call to action: Check CNPG plugin

- Call the help for the CNPG Plugin, run:
`kubectl-cnp help`



Use case

Operator installation



Operator Installation demonstration

NOT NEEDED DURING WORKSHOP
For illustrative purposes.

- Install the operator
- Check the installed CNP Operator in the console
- Discover the features of the Operator in the OpenShift environment
- Check the installed CNP Operator in the web terminal



Install the CNPG Operator and check the it in the web terminal

NOT NEEDED DURING WORKSHOP
For illustrative purposes.

- In the web terminal install the operator:

```
./02_install_operator.sh
```

-> will require admin privs on Openshift



Call to action: Check the it in the web terminal


- In the web terminal check the installation of the operator:

```
./03_check_operator_installed.sh
```



Call to action: Check the installed CNPG Operator in the Openshift console

- In the OpenShift console navigate to:
 - -> Operators
 - -> Installed Operators
 - -> Klick on the Operator installed:



The screenshot shows the Red Hat OpenShift console interface. On the left is a sidebar with navigation options: Home, Favorites, Operators (expanded), OperatorHub, Installed Operators (selected), Helm, Workloads, Networking, Storage, Builds, and Observe. The main content area is titled 'Installed Operators' and shows a table of installed operators. A green arrow points to the 'EDB Postgres for Kubernetes' operator.

Name	Namespace	Managed Namespaces	Status	Last updated	Provided APIs
 cert-manager Operator for Red Hat OpenShift 1.17.0 provided by Red Hat	 cert-manager-operator	 cert-manager-operator	✓ Succeeded Up to date	14. Okt. 2025, 12:19	CertificateRequest Certificate Challenge ClusterIssuer View 4 more...
 EDB Postgres for Kubernetes 1.27.0 provided by EDB	 openshift-operators	All Namespaces	✓ Succeeded Up to date	15. Okt. 2025, 05:43	Backups Cluster Image Catalog Cluster Postgres Database View 6 more...

Call to action: Discover the features of the Operator in the OpenShift environment

Project: openshift-operators ▾

[Installed Operators](#) > Operator details



EDB Postgres for Kubernetes

1.27.0 provided by EDB



Details

YAML

Subscription

Events

All instances

Backups

Cluster Image Catalog

Cluster

Postgres Database

Failover Quorum

Provided APIs

B Backups

PostgreSQL backup (physical base backup)

[+ Create instance](#)

CIC Cluster Image Catalog

A cluster-wide catalog of PostgreSQL operand images

[+ Create instance](#)

C Cluster

PostgreSQL cluster (primary/standby architecture)

[+ Create instance](#)

D Postgres Database

Declarative creation and management of a database on a Cluster

[+ Create instance](#)

FQ Failover Quorum

FailoverQuorum contains the information about the current failover quorum status of a PG cluster

[+ Create instance](#)

IC Image Catalog

A catalog of PostgreSQL operand images

[+ Create instance](#)

P Pooler

P Postgres Publication

SB Scheduled Backups



Use case


Database Migration / Major Version Upgrade




Database Migration

- In this step we will migrate the app database from our existing cluster to the new cluster
- We will create the yaml file with the setting “**import**” in the bootstrap section
- The operator uses internally postgres tools pg_dump and pg_restore
- This method can be used to **migrate** another database or to run **out-of-the-place upgrade**
- Possible settings:

Microservices:



Monolith:



Call to action: Run migration or out of the place upgrade

- In the web terminal 1:

- Run the script:

- ```
./20_upgrade_major_version.sh
```

- Check the cluster creation process:

- ```
kubectl get pods -w
```

- Check the table test in the cluster-user<X>-17, run the command:

- Connect to the cluster:

- ```
oc cnp psql cluster-user<X>-17
```

- Connect to the database app:

- ```
\c app
```


- Run sql commands:

- ```
select version();
```

- ```
select count(*) from test;
```



Grafana Dashboard



Agenda

Start	End	Session
12:30	13:30	Registration & Welcome Lunch
13:30	13:45	Campocamp Vorstellung
13:45	14:00	Red Hat OpenShift & EDB Partnerschaft
14:00	14:15	Einführung in den Postgres-Marketplace
14:15	14:30	Coffee
14:30	15:00	CNPG Operator Reference Architecture
15:00	15:30	Functionalities for CNPG
15:30	17:15	Interactive session & demo
17:15	18:30	Drinks & Aperero

